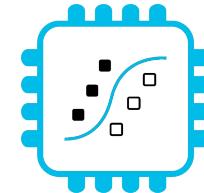
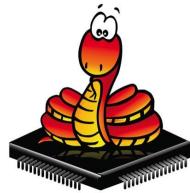
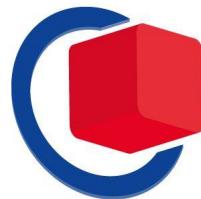


Bridging the TinyML language gap

with MicroPython and emlearn



<https://github.com/emlearn/emlearn>



embeddedworld

Exhibition & Conference

Jon Nordby
jononor@gmail.com
March 2025, Nuremberg



Condition Monitoring

Devices overview

Search Filter by Organization Hotel Manana

Last week Analytics ⓘ

Tech_Storage

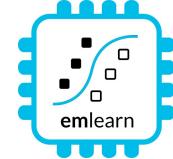
Location	Room	Status
Sandstuvein...	Island 04	Green
Storgata 25...	TBJ 291	Yellow
Chr. Krohgs...	Tech 275	Red
Møllergata 12...	Ohio 3	Green
Ruseløkkveien...	HKM 261	Green
Kristian IVs...	Freeway 273	Yellow
C. J. Hambro...	Oxaca2	Green
Akersgata 65...	Storage_3	Green

Tech 275

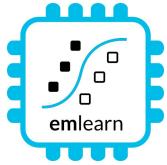
The chart shows a fluctuating line representing sensor data over a week. A red shaded rectangular area highlights a period of high values between Wednesday and Thursday, indicating a specific event or alert period.

Trusted by Nordic market leaders

Outline / agenda

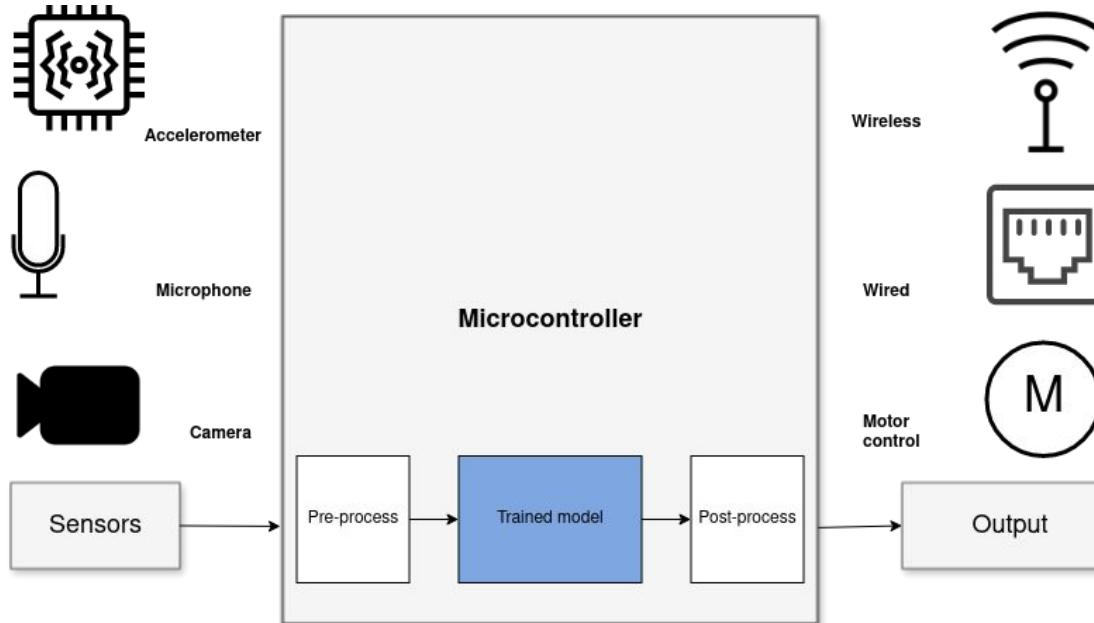


1. Brief introduction to **TinyML**
2. TinyML with emlearn
 - **emlearn** - overview, C library
 - emlearn + **Zephyr RTOS**
2. Bridging the TinyML language gap
 - **MicroPython**
 - emlearn-micropython



TinyML overview

Sensor data analysis with Machine Learning



Benefits vs sending data to cloud

- Stand-alone
- Low latency
- Power efficiency
- Privacy
- Low cost

1) Read sensors → 2) Process data * → 3) Transmit/act on data

* including Digital Signal Processing (DSP) and Machine Learning (ML)

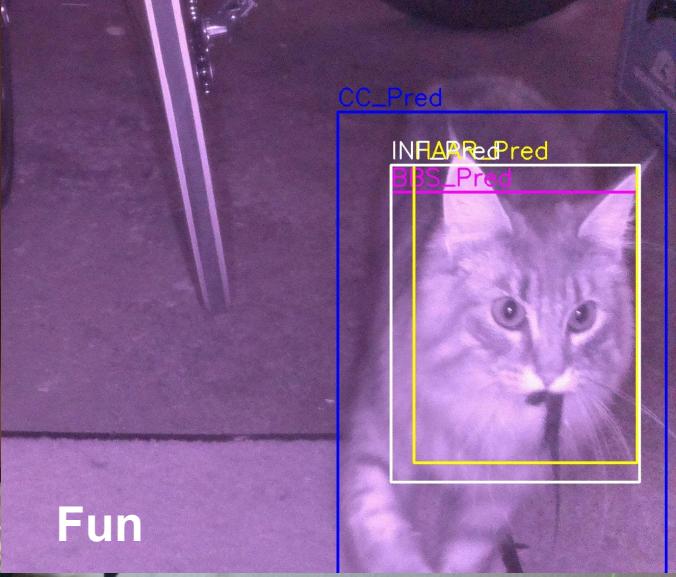
Hey Google...



Consumer Tech



Industrial



Fun



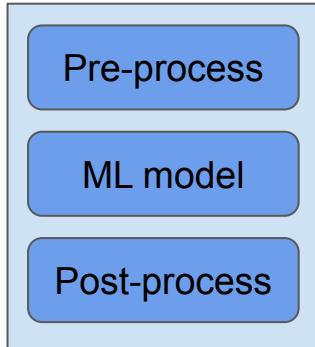
TinyML challenge: PC/host vs embedded/device pipeline



1) Train model
on PC



RAM (min)
10 GB+

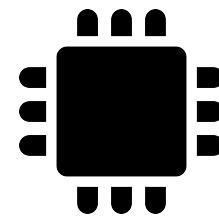


“Data Scientist land”

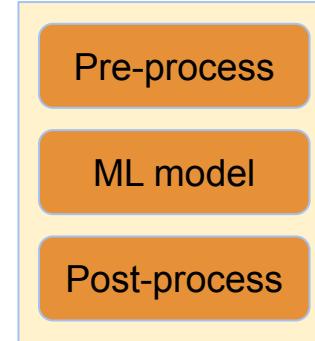
2) Run
on Device



Sensor



RAM (max):
1 MB

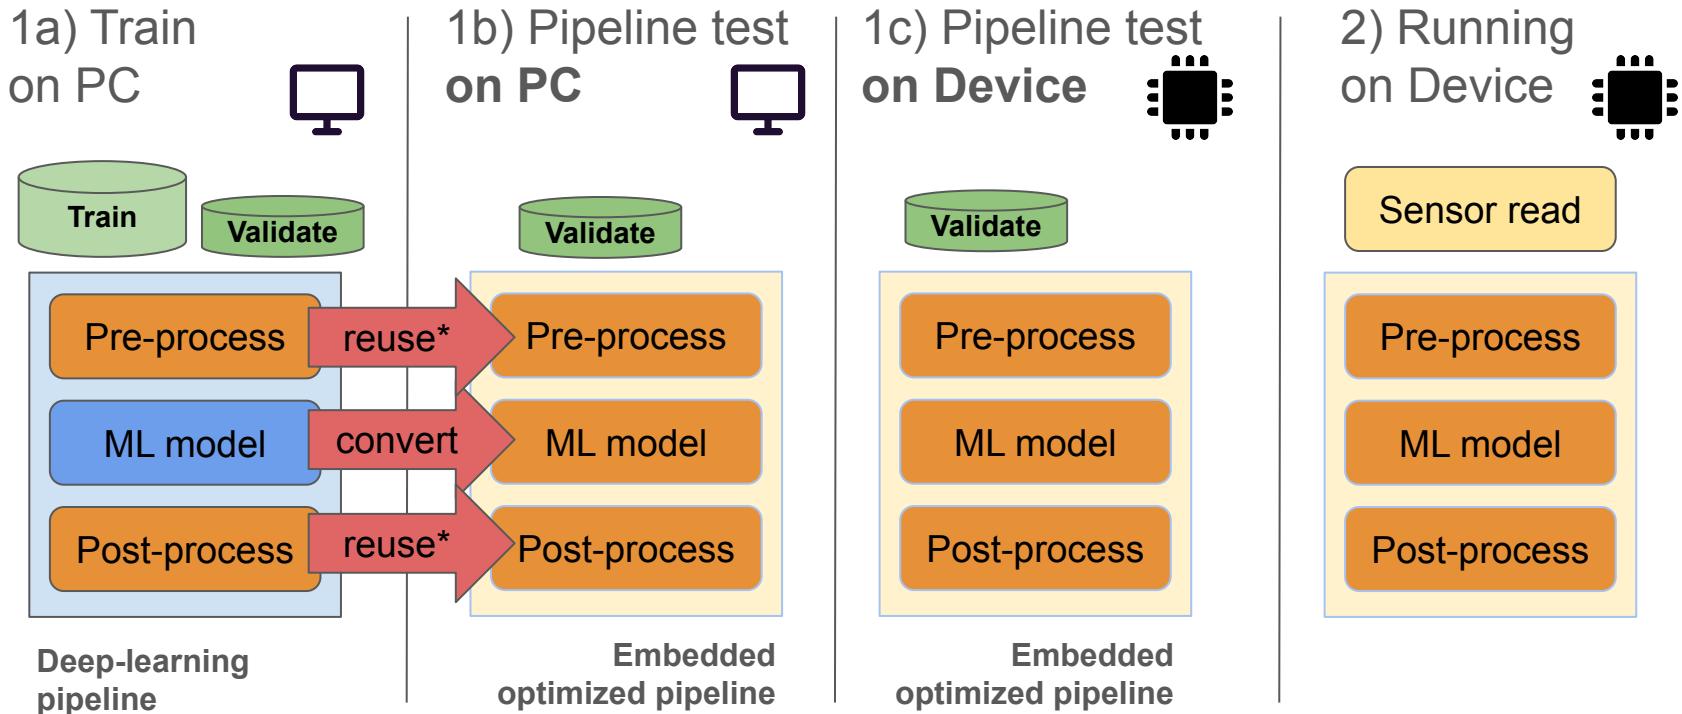


“Firmware Engineer land”

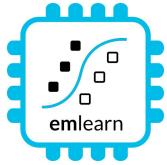
Challenges

- Very different constraints in the two environments
- The two pipelines must be 100% compatible
- Pre- and post-processing often application-specific, Turing complete
- Different languages, tools, libraries, cultures

Strategy: Step-wise validation, code-sharing



* Requirements: Pre- and -post processing must be portable between device and PC
Can be done in embedded C/C++ etc, or (MicroPython)



About emlearn

Started in 2018

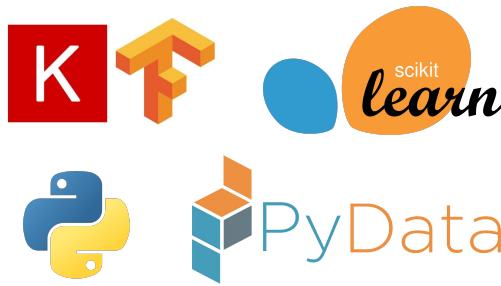
Open-source (MIT)

1. Train on PC

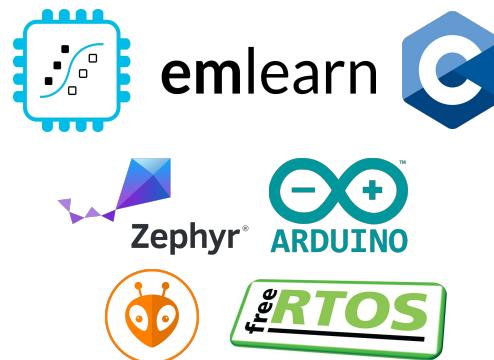
\$ pip install emlearn

Convenient training

- Model creation in Python
- Use standard libraries
 - a. scikit-learn
 - b. Keras
- One-line to **export to device**
- Verification tools included

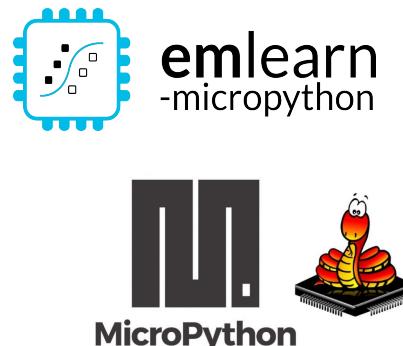


2. Deploy on device



Embedded Friendly

- Portable C99 code
- No dynamic allocations
- Header-only
- High test coverage
- Integer/fixed-point math *
- Small. 2 kB+ FLASH

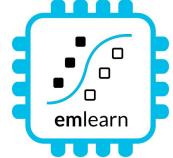


Convenient & Efficient

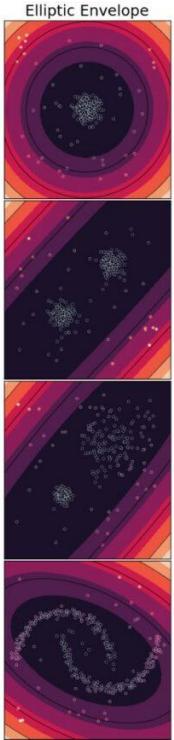
- Portable MicroPython
- Single .mpy file
- No dynamic allocations
- High test coverage
- Integer/fixed-point math *
- Small. 2 kB+ FLASH

All code runs both on host/PC and device

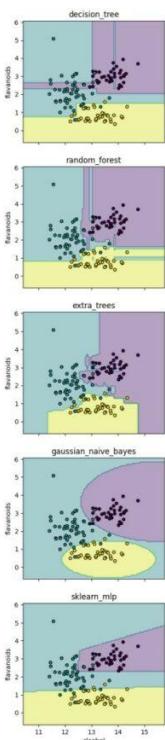
Supported tasks



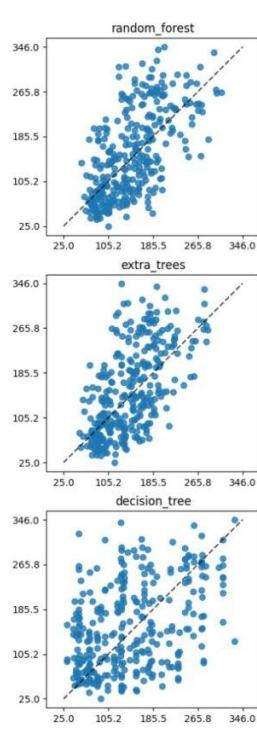
Anomaly Detection



Classification



Regression



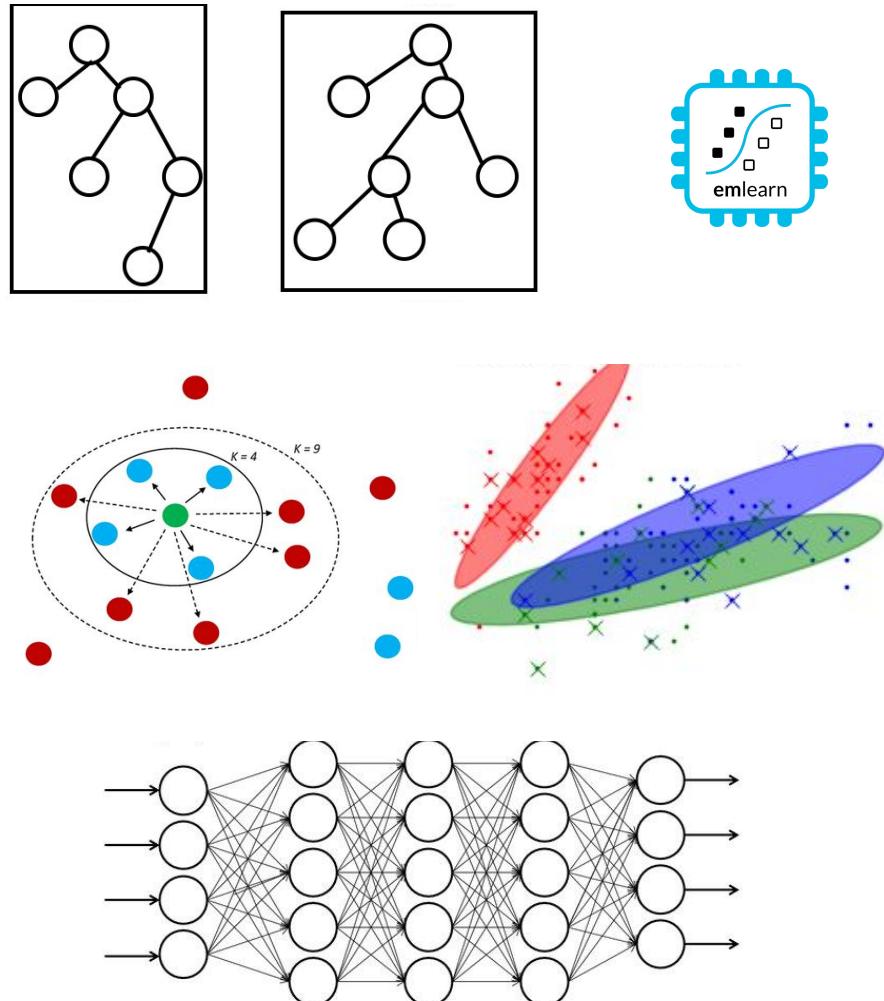
Supports the most common tasks
for embedded
& sensor data use cases.

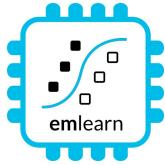
- Classification
- Regression
- Anomaly Detection

Supported models

Selection of simple & effective
embedded-friendly models

- Decision Trees (DT)
- Random Forest (RF)
- K Nearest Neighbors (KNN)
- Gaussian Mixture Models (GMM)
- Multi-Layer-Perceptron (MLP)
- Convolutional Neural Network (CNN)





Projects using emlearn

Many real-world uses
across the globe

Referred in 40+ publications

Health monitoring for cattle

Accelerometer data

used to detect abnormal behavior.

Can indicate **health issues** or other problems

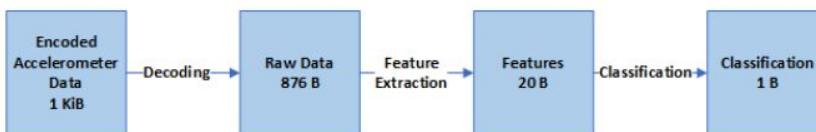
Classified using a **Decision Tree**

lying/walking/standing/grazing/ruminating

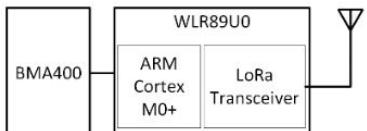
Activity is transmitted using **LoRaWAN**

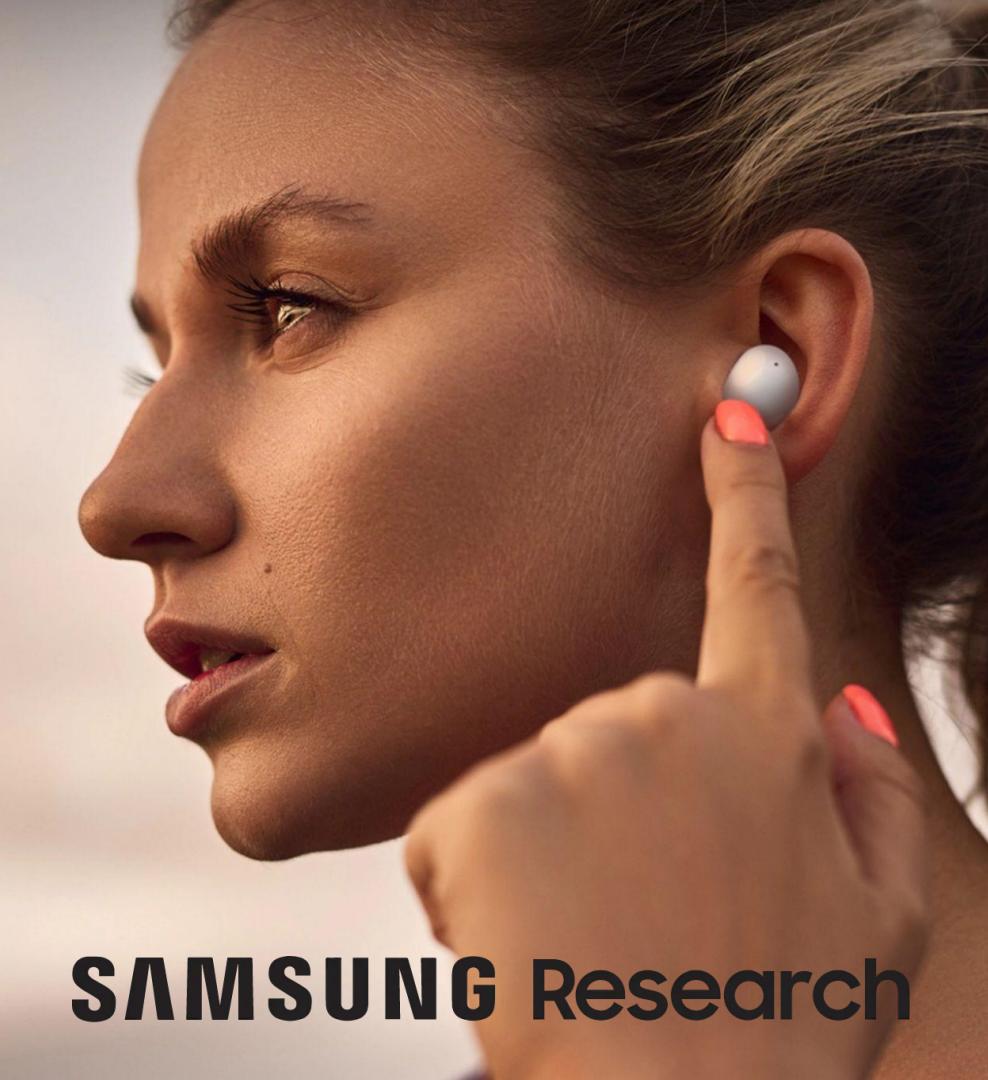
Running ML on-sensor: **under 1 mW**

50 times lower power than sending raw data



Power Efficient Wireless Sensor Node through Edge Intelligence
Abhishek Damle (2022)





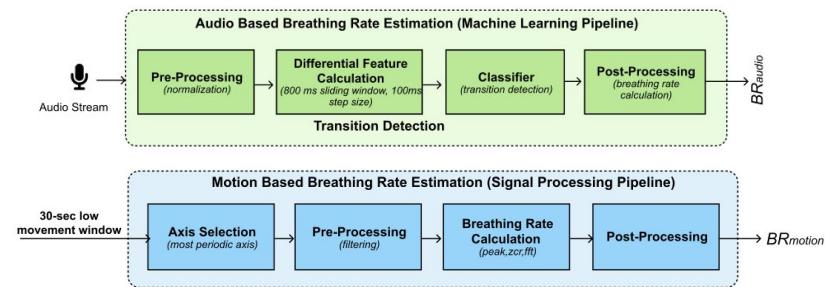
SAMSUNG Research

Health monitoring earable

Breathing rate is critical to monitor for persons with **respiratory health problems**.

Monitoring **for every-day use**, not just in clinical context. Using earable to replace specialized devices.

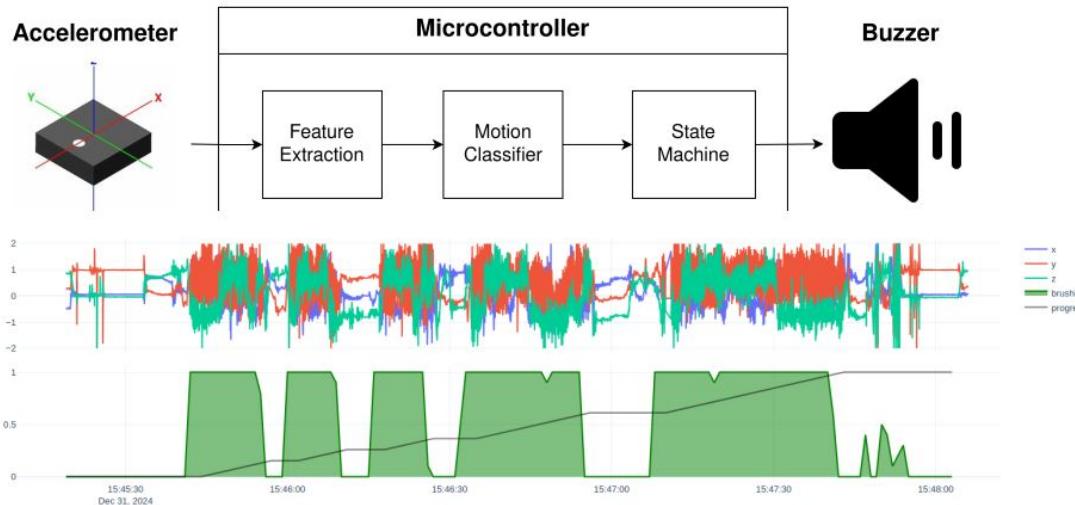
Random Forest classifier for audio.



*Remote Breathing Rate Tracking in Stationary Position
Using the Motion and Acoustic Sensors of Earables
Tousif Ahmed et.al (2023)*

Automatic toothbrush timer

<https://github.com/jonnor/toothbrush/>



- Stock hardware: M5StickC PLUS 2 by M5Stack
- ~500 lines of MicroPython code
- Collected and labeled data in ~1 hour

Data collection and training tools:

https://github.com/emlearn/emlearn-micropython/tree/master/examples/har_trees



Brushing time
00:31 / 02:00
26 % done



emlearn for C

Quick walkthrough

Train and export a model



1. Train using standard Python ML libraries.

```
from keras import ...      A) keras neural  
                           network  
  
model = Sequential([  
    Dense(16, input_dim=n_features, activation='relu'),  
    Dense(8, activation='relu'),  
    Dense(1, activation='sigmoid'),  
])  
model.compile(...)  
model.fit(X_train, Y_train, epochs=1, batch_size=10)
```

```
from sklearn.neural_network import MLPClassifier  
model = MLPClassifier(hidden_layer_sizes=(100,50,25))  
  
model.fit(X_train, Y_train)  
  
B) scikit-learn neural network
```

2. Use `emlearn.convert()` and `.save()`

```
import emlearn  
  
cmodel = emlearn.convert(model, method='inline')  
  
cmodel.save(file=mynet_model.h', name='mynet')
```



```
#include <eml_net.h>  
  
static const float mynet_layer_0_biases[8] = { -0.015587f, -0.005395f, -0.010957f, 0.015883f ....  
static const float mynet_layer_0_weights[24] = { -0.256981f, 0.041887f, 0.063659f, 0.011013f, ....  
static const float mynet_layer_1_biases[4] = { 0.001242f, 0.010440f, -0.005309f, -0.006540f };  
static const float mynet_layer_1_weights[32] = { -0.577215f, -0.674633f, -0.376140f, 0.646900f, ....  
static float mynet_buf1[8];  
static float mynet_buf2[8];  
static const EmlNetLayer mynet_layers[2] = {  
    { 8, 3, mynet_layer_0_weights, mynet_layer_0_biases, EmlNetActivationRelu },  
    { 4, 8, mynet_layer_1_weights, mynet_layer_1_biases, EmlNetActivationSoftmax }  
};  
static EmlNet mynet = { 2, mynet_layers, mynet_buf1, mynet_buf2, 8 };  
  
int32_t  
mynet_predict(const float *features, int32_t n_features)  
{  
    return eml_net_predict(&mynet, features, n_features);  
}  
.....
```

Example of generated code

Using the C code



3. #include and call predict()

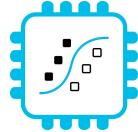
```
// Include the generated model code  
#include "mynet_model.h"  
  
// index for the class we are  
detecting  
#define MYNET_VOICE 1  
  
// Buffers for input data  
#define N_FEATURES 6  
float features[N_FEATURES];  
  
#define DATA_LENGTH 128  
int16_t  
sensor_data[DATA_LENGTH];
```

setup

```
// Get data and pre-process it  
read_microphone(sensor_data, DATA_LENGTH);  
preprocess_data(sensor_data, features);  
  
// Run the model  
out = mynet_predict(features, N_FEATURES);  
  
// Do something with results  
if (out == MYNET_VOICE) {  
    set_display("voice detected");  
} else {  
    set_display("");  
}
```

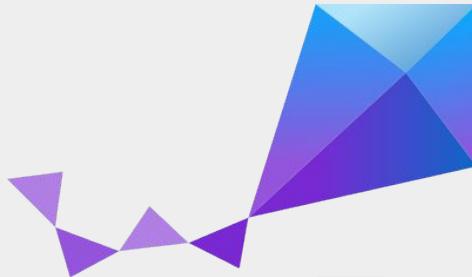
loop





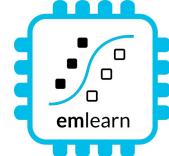
emlearn

emlearn
on Zephyr



Zephyr®

emlearn module for Zephyr



1. Add project to west.yml

```
projects:  
  #... existing projects  
  - name: emlearn  
    url: https://github.com/emlearn/emlearn  
    revision: master
```

2. Enable in your project

```
CONFIG_EMLEARN=y
```

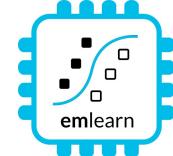
3. Use in your code

```
#include <eml_net.h>
```

First class support
since February 2024

[https://emlearn.readthedocs.io
/en/latest/
getting_started_zephyr.html](https://emlearn.readthedocs.io/en/latest/getting_started_zephyr.html)

Zephyr sensor data APIs

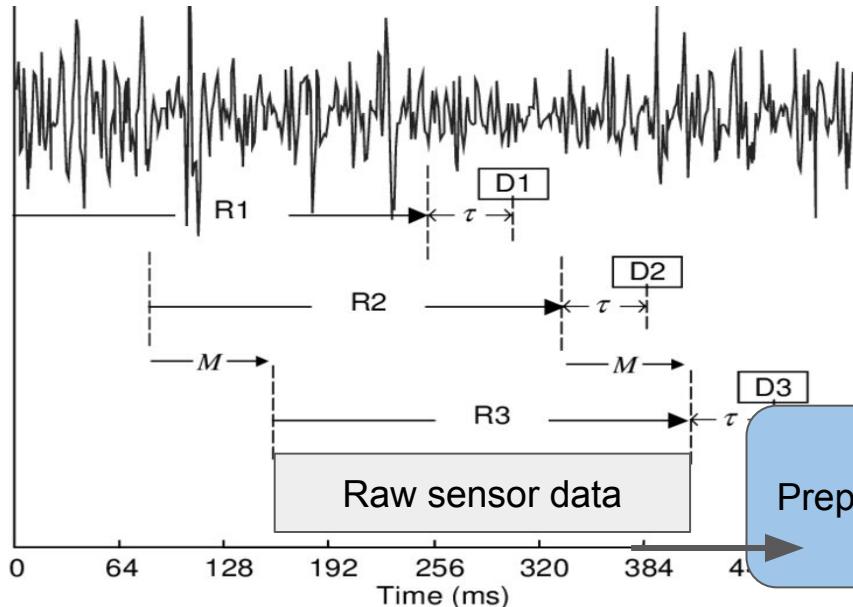
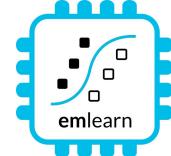


- ★ Unified hardware API across all platforms
- ★ Large amount of drivers included out-of-the-box

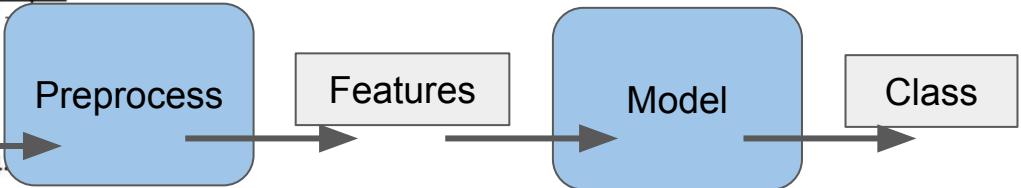
Examples

IMU (accelerometer/gyro)	Sensors API	https://docs.zephyrproject.org/latest/hardware/peripherals/sensor/index.html
Audio	Digital mic (DMIC) Digital Audio Interface (DAI)	https://docs.zephyrproject.org/latest/hardware/peripherals/audio/dmic.html
Images (camera)	Video API	https://docs.zephyrproject.org/latest/hardware/peripherals/video.html

ML on streams: Continuous classification



The sensor data stream
is sliced into overlapping windows.
Each window processed independently

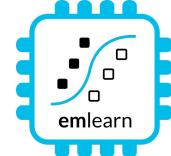


[x1 y1 z1
x2 y2 z2
....]

[42, 4002, ..., 329]

“Jumping Jacks”

Chunked sensor reader for Zephyr



```
struct sensor_chunk_msg chunk;

if (!device_is_ready(sensor)) {
    printk("sensor: device %s not ready.\n", sensors[i]->name);
    return 0;
}

// Setup sensor
set_sampling_freq(sensor, SAMPLINGRATE);

// Start high-priority thread for collecting data
sensor_chunk_reader_start(&sensor_reader);

while (1) {

    // check for new data
    const int get_status = k_msqq_get(sensor_reader.queue, &chunk, K_NO_WAIT);
    if (get_status != 0) {

        process(chunk.buffer, chunk.length);
    }

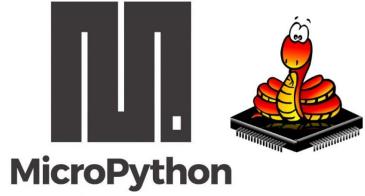
    k_msleep(100);
}
```

Work-in-progress

Practical examples using sensor data for Zephyr
<https://github.com/emlearn/emlearn/issues/108>

Dedicated thread handling time-sensitive fetching

Queue with complete chunks of data
Allows application to process at convenient time



Bridging the TinyML language gap

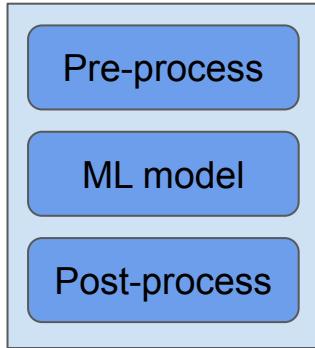
TinyML challenge: PC/host vs embedded/device pipeline



1) Train model
on PC

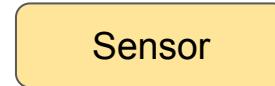


RAM (min)
10 GB+

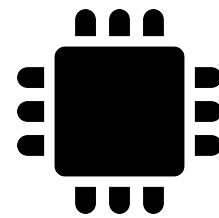


“Data Scientist land”

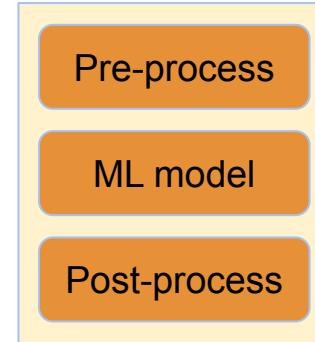
2) Run
on Device



Sensor



RAM (max):
1 MB



Challenges

- Very different constraints in the two environments
- The two pipelines must be 100% compatible
- Pre- and post-processing often application-specific, Turing complete
- Different languages, tools, libraries, cultures

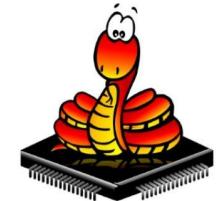
“Firmware Engineer land”

MicroPython - Introduction

Implements a subset of Python 3.x

For devices with 16 kB+ RAM

Supports 8+ microcontroller families



As compatible with CPython as possible, within constraints.

No CFFI or C module compatibility!

TLDR: Small .py scripts will mostly work with minor mods

But not standard PyData stack

“mip” package manager

Can load C modules at runtime!

More info: <https://micropython.org>

RP2040



C modules *

Defines a Python module with API.
functions/classes etc.

**Implemented by users,
libraries, or be part of
MicroPython core.**

Can be **portable** or
hardware/platform specific

* Or other language which
compiles to C, or exposes C API

[https://github.com/
vshymanskyy/wasm2mpy](https://github.com/vshymanskyy/wasm2mpy)

C++, Rust, Zig, TinyGo, TypeScript

```
// Include the header file to get access to the MicroPython API
#include "py/dynruntime.h"

// Helper C function to compute factorial
static mp_int_t factorial_helper(mp_int_t x) {
    if (x == 0) {
        return 1;
    }
    return x * factorial_helper(x - 1);
}

// DEFINE FUNCTION. Callable from Python
static mp_obj_t factorial(mp_obj_t x_obj) {
    mp_int_t x = mp_obj_get_int(x_obj);
    mp_int_t result = factorial_helper(x);
    return mp_obj_new_int(result);
}
static MP_DEFINE_CONST_FUN_OBJ_1(factorial_obj, factorial);

// MODULE ENTRY
mp_obj_t mpy_init(mp_obj_fun_bc_t *self, size_t n_args, size_t n_kw, mp_obj_t *args) {
    // Must be first, it sets up the globals dict and other things
    MP_DYNRUNTIME_INIT_ENTRY

    // Register function in the module's namespace
    mp_store_global(MP_QSTR_factorial, MP_OBJ_FROM_PTR(&factorial_obj));

    // This must be last, it restores the globals dict
    MP_DYNRUNTIME_INIT_EXIT
}
```

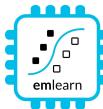


Native module (.mpy) VS External C module

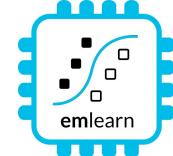
	Native module	External C module
Installable at runtime	Yes, as .mpy file	No. Must be included in firmware image
Requires SDK/toolchain	No (only to build)	Yes
Code executes from	RAM	FLASH
Limitations	No libc / libm linked * No static BSS *	None
Maturity	Low *	Excellent
Documentation	https://docs.micropython.org/en/latest/develop/natmod.html	https://docs.micropython.org/en/latest/develop/cmodules.html

* Improved greatly in upcoming MicroPython (1.25+).

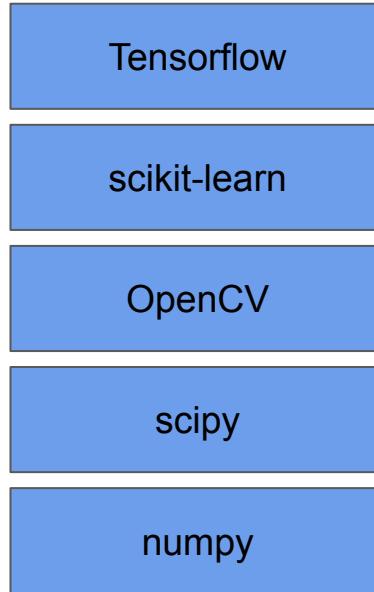
Contributions by Volodymyr Shymanskyy, Alessandro Gatti, Damien George, and others



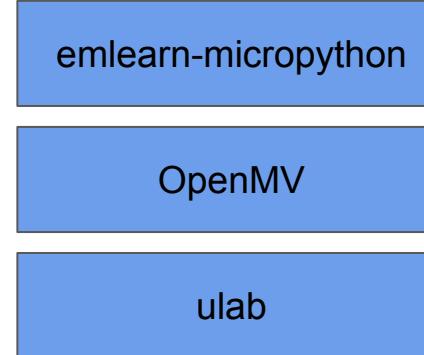
MicroPython Data Science ecosystem



C_{Python}



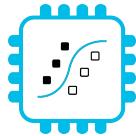
MicroPython



<https://openmv.io/>

<https://github.com/v923z/micropython-ulab>

Libraries existing that implement a lot of embedded-relevant functionality from PyData stack.



emlearn
-micropython

emlearn for MicroPython

Python is the lingua franca
for Machine Learning

How can we enable people to build
TinyML solution also using Python?

MicroPython library for emlearn

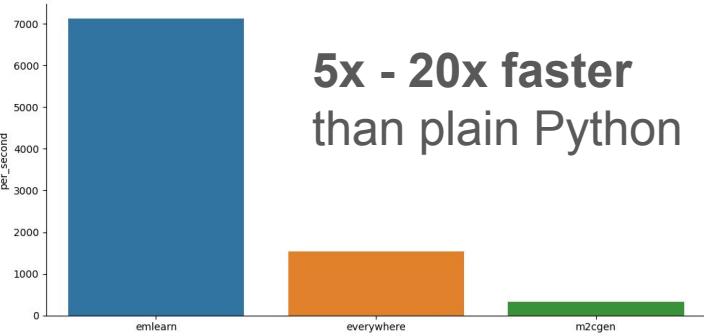


Combine the **convenience, familiarity and productivity of Python** - with the **speed of C**

Modules installable at runtime (.mpy)

emlearn_iir
emlearn_trees
emlearn_fft
emlearn_cnn
emlearn_neighbors

Infinite Impulse Response filters
Random Forest
Fast Fourier Transform
Convolutional Neural Networks
K-nearest Neighbors



5x - 20x faster
than plain Python

Image classification+
On-device learning

<https://github.com/emlearn/emlearn-micropython/>

Install & Export model



```
import emlearn
```

Export model on PC

```
estimator = train_model()  
converted = emlearn.convert(estimator)  
converted.save(name='gesture', format='csv', file='gesture_model.csv')
```

Copy model to device

```
mpremote fs cp gesture_model.csv :
```

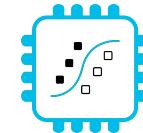
Prebuilt binary

No need to setup
C toolchain and SDK

Copy library to device

```
mpremote mip install https://emlearn.github.io/....../xtensawin_6.2/emltrees.mpy
```

Load model & run



emlearn
-micropython

```
import emltrees
# Instantiate model. Capacity for trees, decision nodes, classes
model = emltrees.new(20, 200, 10)

# Load model "weights"
with open('gesture_model.csv') as f:
    emltrees.load_model(model, f)
```

Load model on device

```
# Read sensor data
sensor.fifo_read(samples)

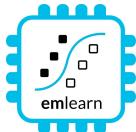
# Preprocess features
processor.run(samples, features)

# Run model
out = model.predict(features)
```

Run inference

Activity tracker

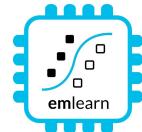
Accelerometer



Random Forest
classifier
emlearn_trees

Noise monitor

Microphone



Infinite Impulse
Response filters
emlearn_iir

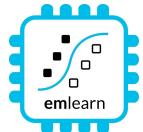
Image Classifier

Camera

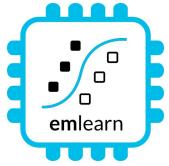


Examples at
[https://github.com/
emlearn/emlearn-micropython
#examples](https://github.com/emlearn/emlearn-micropython#examples)

```
if is_MyCat(img):  
    open_door()
```

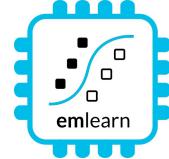


Convolutional
Neural Network
emlearn_cnn



Outro

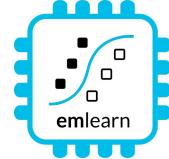
Summary 1/2



TinyML pipelines must be
compatible between train/PC and deploy/device

Recommend using **portable code** with
automated validation tests on both device and PC

Summary 2/2



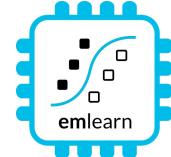
emlearn is a portable open-source project
to help **deploy ML pipelines to microcontrollers**

emlearn supports both **C** and **MicroPython**

MicroPython with C modules enables writing
TinyML applications primarily in Python

More resources

Official documentation



<https://emlearn-micropython.readthedocs.io>

<https://emlearn.readthedocs.io>

PyCon Berlin 2024: **Machine Learning on microcontrollers using MicroPython and emlearn**

<https://www.youtube.com/watch?v=S3GjLr0ZIEQ>

TinyML EMEA 2024: **emlearn - scikit-learn for microcontrollers and embedded systems**

<https://www.youtube.com/watch?v=LyO5k1VMdOQ>

FOSDEM 2025: **MicroPython - Python for microcontrollers and Embedded Linux**

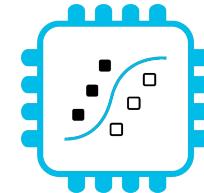
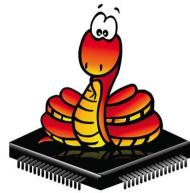
<https://fosdem.org/2025/schedule/event/fosdem-2025-4525-micropython-python-for-microcontrollers-and-embedded-linux/>

FOSDEM 2025: **Milliwatt sized Machine Learning on microcontrollers with emlearn**

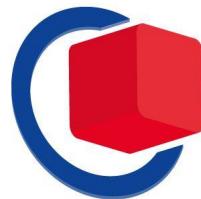
<https://fosdem.org/2025/schedule/event/fosdem-2025-4524-milliwatt-sized-machine-learning-on-microcontrollers-with-emlearn/>

Bridging the TinyML language gap

with MicroPython and emlearn



<https://github.com/emlearn/emlearn>



embeddedworld

Exhibition & Conference

Jon Nordby
jononor@gmail.com
March 2025, Nuremberg

Bonus

1 dollar TinyML

<https://hackaday.io/project/194511-1-dollar-tinyml>

Research question:

Can we make a useful TinyML system
<1 USD in total component cost (BOM)?

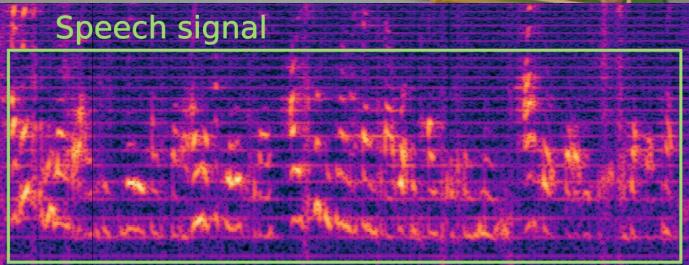
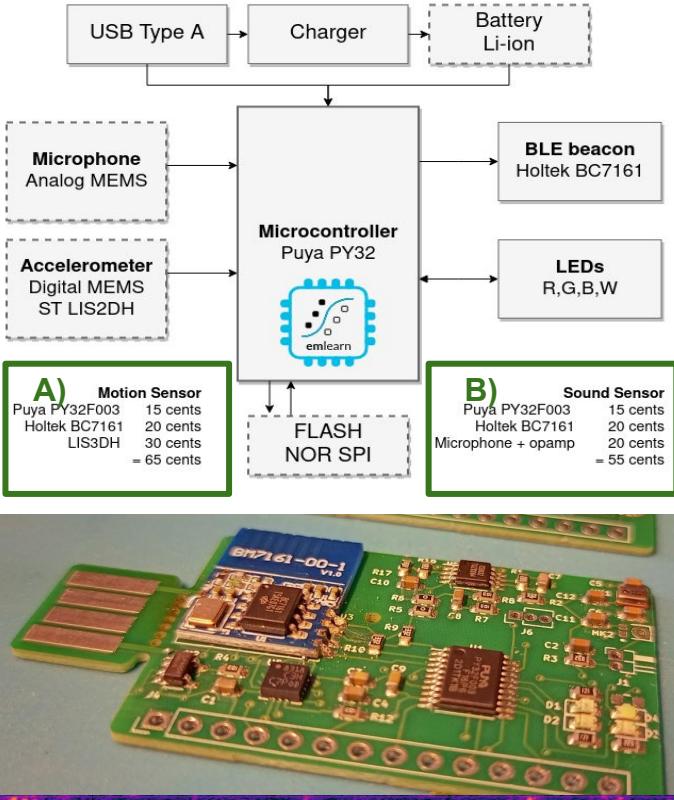
Preliminary answer: YES!

Either motion (accelerometer)
or sound sensor (microphone)

Constraints: 8 kB RAM / 64 kB FLASH

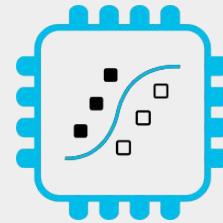
Current status:

Audio streaming works
Revision 2 ordered



Sensor nodes

with MicroPython
and emlearn



emlearn
-micropython

[https://github.com/
emlearn/emlearn-micropython](https://github.com/emlearn/emlearn-micropython)

Task feasibility versus microcontroller type

	RAM (kB)	FLASH (kB)	CPU (CoreMark)		IMU 100 hz 5 s	Sound 16 kHz 1 s	Image 64x64 px 1 fps
Arduino Uno ATMega 328 @ 16 Mhz	2	32	4.0		✓	✗	✗
Arduino Nano BLE 33 NRF52840 ARM Cortex M4F @ 64 Mhz	256	1 000	215		✓	✓	✗
Arduino Nano ESP32 ESP32-S3 @ 240 Mhz	8 000	16 000	613 (per core)		✓	✓	✓
Teensy 4.0 ARM Cortex M7 @ 600 Mhz	1 000	2 000	2313 (per core)		✓	✓	✓

Audio input - machine.I2S

Digital microphone or external audio ADC

Can be done using I2S protocol

On ports **ESP32, STM32, RP2, NRF52**

PDM protocol not supported :(

Example code

https://github.com/emlearn/emlearn-micropython/tree/master/examples/soundlevel_iir

<https://github.com/miketeachman/micropython-i2s-examples>

```
# I2S audio input
from machine import I2S
audio_in = I2S(0, sck=Pin(26), ws=Pin(32), sd=Pin(33),
               mode=I2S.RX, bits=16, format=I2S.MONO, rate=16000,
               )

# allocate sample arrays
chunk_samples = int(AUDIO_SAMPLERATE * 0.125)
mic_samples = array.array('h', (0 for _ in range(chunk_samples))) # int16
# memoryview used to reduce heap allocation in while loop
mic_samples_mv = memoryview(mic_samples)
# global to share state between callback and main
soundlevel_db = 0.0

meter = SoundlevelMeter(buffer_size=chunk_samples, samplerate=16000)

def audio_ready_callback(arg):
    # compute soundlevel
    global soundlevel_db
    soundlevel_db = meter.process(mic_samples)
    # re-trigger audio callback
    _ = audio_in.readinto(mic_samples_mv)

def main():
    # Use Non-Blocking I/O with callback
    audio_in.irq(audio_ready_callback)
    # Trigger first audio readout
    audio_in.readinto(mic_samples_mv)

    while True:
        render_display(db=soundlevel_db)
        time.sleep_ms(200)

if __name__ == '__main__':
    main()
```

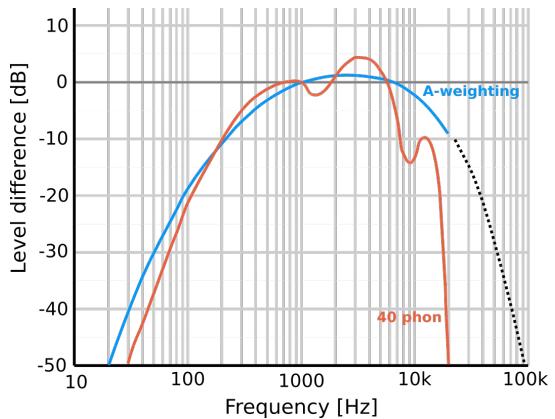
Noise sensor using emlearn_iir

Using `machine.I2S`
16 kHz samplerate

A weighting implemented with
Infinite Impulse Response (IIR) filter

`emlearn_iir`
25% CPU usage total

pure MicroPython
900% CPU - not feasible



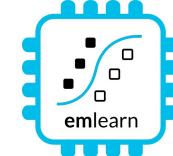
```
# 6th order filter. 3x Second Order Sections "biquad"
a_filter_16k = [
    1.0383002230320646, 0.0, 0.0, 1.0, -0.016647242439959593, 6.928267
    1.0, -2.0, 1.0, 1.0, -1.7070508390293027, 0.7174637059318595,
    1.0, -2.0, 1.0, 1.0, -1.9838868447331497, 0.9839517531763131
]
self.frequency_filter = IIRFilter(a_filter_16k)
...
self.frequency_filter.process(samples)
```

Complete example code

https://github.com/emlearn/emlearn-micropython/tree/master/examples/soundlevel_iir



Human Activity Detection with emlearn_trees



Using Random Forest classifier trained with scikit-learn

```
import emltrees
model = emltrees.new(10, 1000, 10)
with open('eml_digits.csv', 'r') as f:
    emltrees.load_model(model, f)

features = array.array('h', ...)
out = model.predict(features)
```

Performance comparison

10 trees, max 100 leaf nodes, “digits” dataset

	Inference time	Program space
m2cgen	60.1 ms	179 kB
everywhereml	17.7 ms	154 kB
emlearn	1.3 ms	15 kB



https://github.com/emlearn/emlearn-micropython/tree/master/examples/har_trees

emlearn is **10x faster and 10x more space efficient** compared to generating Python code

Activity Tracker - Feature Extraction

Statistical summarizations are useful time-series features, sufficient for basic Human Activity Recognition.

! Preprocessing *must be compatible*
between training on host PC (CPython) and device (MicroPython)

Solution: Write preprocessor for MicroPython, re-use in Python

```
subprocess('micropython preprocess.py data.npy features.npy')
```

Alternative: (when using common MicroPython/CPython subset)

```
import mypreprocessor.py
```

Using micropython-npyfile to read/write Numpy .npy files
<https://github.com/jonnor/micropython-npyfile/>

```
l = sorted(list(v))
l2 = [x*x for x in l]
sm = sum(l)
sqsum = sum(l2)
avg = sum(l) / len(l)

median = l[MEDIAN]
q25 = l[Q1]
q75 = l[Q3]
iqr = (l[Q3] - l[Q1])

energy = ((sqsum / len(l2)) ** 0.5)
std = ((sqsum - avg * avg) ** 0.5)
```

[https://github.com/emlearn/emlearn-micropython
/blob/master/examples/har_trees/timebased.py](https://github.com/emlearn/emlearn-micropython/blob/master/examples/har_trees/timebased.py)

Time-based features extraction
Are Microcontrollers Ready for Deep
Learning-Based Human Activity Recognition?
Atis Elsts, and Ryan McConville
<https://www.mdpi.com/2079-9292/10/21/2640>

Training model on dataset

Using a scikit-learn based pipeline.

```
# Setup subject-based cross validation
splitter = GroupShuffleSplit(n_splits=n_splits, test_size=0.25,
                             random_state=random_state)

# Random Forest classifier
clf = RandomForestClassifier(random_state = random_state,
                            n_jobs=1, class_weight = "balanced")

# Hyper-parameter search
search = GridSearchCV(clf, param_grid=hyperparameters,
                      scoring=metric, refit=metric, cv=splitter)
search.fit(X, Y, groups=groups)
```

```
(venv) [jon@jon-thinkpad har_trees]$ MIN_SAMPLES_LEAF=150,200,400 python har_train.py
-dataset har_exercise_1 --window-length 400 --window-hop 10
2024-12-04 12:54:52 [info] data-loaded dataset=har_exercise_1
duration=0.016095876693725586 samples=32000
2024-12-04 12:54:56 [info] feature-extraction-done dataset=har_exercise_1
duration=4.534412145614624 labeled_instances=1952 total_instances=1952
Class distribution
activity
jumpingjack    549
lunge           488
other            488
squat            427
Name: count, dtype: int64
Model written to ./har_exercise_1_trees.csv
Testdata written to ./har_exercise_1.testdata.npz
Results
   n_estimators  min_samples_leaf  mean_train_f1_micro  mean_test_f1_micro
0             10              150          0.996311        0.962705
1             10              200          0.995628        0.956557
2             10              400          0.986202        0.920902
```

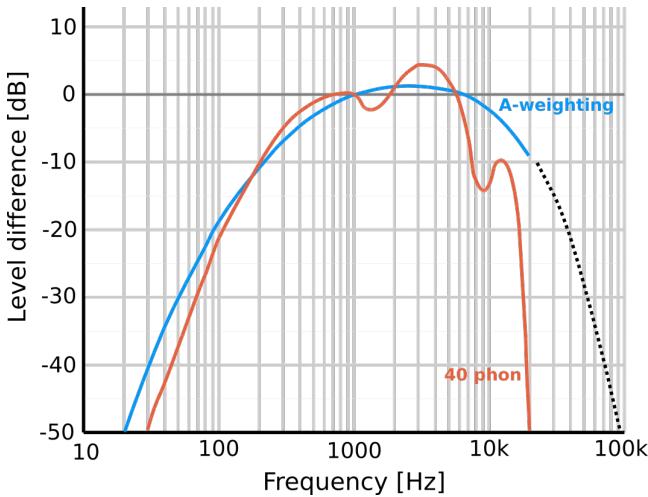
har_train.py

```
import emlearn
converted = emlearn.convert(clf)
converted.save(name='gesture', format='csv', file='model.csv')
```

Sound sensor - IIR filter

Standard sound level measurements are **A-weighted**. To approximate human hearing.

Typically, implemented using **Infinite Impulse Response (IIR) filters**.



```
class IIRFilter():
    def __init__(self, coefficients : array.array):
        stages = len(coefficients)//6
        self.sos = coefficients
        self.state = array.array('f', [0.0]*(2*stages))

    @micropython.native
    def process(self, samples : array.array):
        stages = len(self.sos)//6
        for i in range(len(samples)):
            x = samples[i]
            for s in range(stages):
                b0, b1, b2, a0, a1, a2 = self.sos[s*6:(s*6)+6]
                # compute difference equations of transposed direct form II
                y = b0*x + self.state[(s*2)+0]
                self.state[(s*2)+0] = b1*x - a1*y + self.state[(s*2)+1]
                self.state[(s*2)+1] = b2*x - a2*y
                x = y
            samples[i] = x
```

1100 ms 900% CPU
Native emitter
Float

Too slow by ~10x

```
# 6th order filter. 3x Second Order Sections "biquad"
a_filter_16k = [
    1.0383002230320646, 0.0, 0.0, 1.0, -0.016647242439959593, 6.9282670,
    1.0, -2.0, 1.0, 1.0, -1.7070508390293027, 0.7174637059318595,
    1.0, -2.0, 1.0, 1.0, -1.9838868447331497, 0.9839517531763131
]
self.frequency_filter = IIRFilter(a_filter_16k)
...
self.frequency_filter.process(samples)
```



Sound sensor - IIR filter

Using **emliir.mpy** native module helps a lot.

BUT - conversion from float/int16 too slow
Also needs a native module

IIR filter only
Using emliir.mpy
native module
30 ms 20% CPU - OK

```
import emliir

class IIRFilterEmlearn:

    def __init__(self, coefficients):
        c = array.array('f', coefficients)
        self.iir = emliir.new(c)
    def process(self, samples):
        self.iir.run(samples)
```

```
@micropython.native
def float_to_int16(inp, out):
    for i in range(len(inp)):
        out[i] = int(inp[i]*32768)
```

```
@micropython.native
def int16_to_float(inp, out):
    for i in range(len(inp)):
        out[i] = inp[i]/32768.0
```

```
int16_to_float(samples, self.float_array)
self.frequency_filter.process(self.float_array)
float_to_int16(self.float_array, samples)
```

But need to convert data types
Adds 70ms+ with micropython.native
Too slow! Total > 100 ms
Must create native module



Camera input

Not part of standard APIs yet

micropython-camera-API

Proposed API and implementation
(ESP32 only, for now)

[https://github.com/
cnadler86/micropython-camera-API](https://github.com/cnadler86/micropython-camera-API)

OpenMV

<https://openmv.io/>

Custom MicroPython distribution
Focused on Computer Vision / Machine Vision



Image classification with emlearn_cnn

Convolutional Neural Network (CNN)

Running on ESP32-S3 (example)

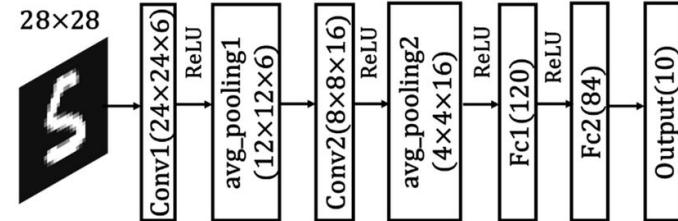
Input dimensions: 32x32 px - 96x96 px

Layers: 3 - 4 layers.

Framerate: 1 - 10 FPS

Complete example code

https://github.com/emlearn/emlearn-micropython/tree/master/examples/mnist_cnn



```
import tinymaix_cnn # from emlearn-micropython

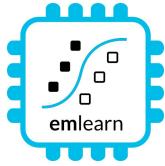
with open('cat_classifier.tmdl', 'rb') as f:
    model_data = array.array('B', f.read())
    model = tinymaix_cnn.new(model_data)

while True:

    raw = read_camera()
    img = preprocess(raw)
    classification = model.predict(img)

    if classification == MY_CAT:
        open_door()

    machine.lightsleep(500)
```



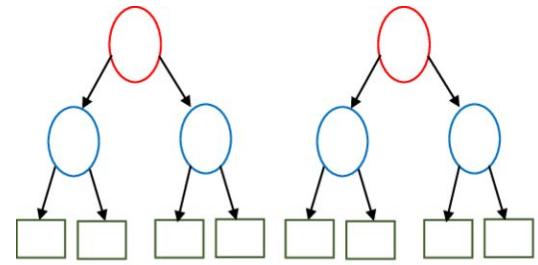
Tree-based ensembles

Random Forest / Decision Tree

The most popular model in emlearn

emlearn trees: Leaf compression

scikit-learn decision trees store class probabilities in leaf-nodes. Size: $n_{\text{leaves}} * n_{\text{classes}} * 4$ bytes
Consequence: Leaves take a large amount of space



emlearn (since 2019) does **hard majority voting** instead
And does **leaf-node de-duplication** across all trees

Size: $n_{\text{classes}} * 1 * 1$ byte
Saving 50-80% of space

Tradeoff: May lead to reduced predictive performance with some leaf values

Lossless	0.0	1.0	0.0	0.0	1
Not ideal	0.0	0.0	1.0	0.0	2
	0.0	0.4	0.6	0.0	2
	0.15	0.20	0.30	0.15	2

.....

emlearn trees: Inference modes

Data structure (“loadable”)

- + Can be loaded at runtime
- Only supports float

```
static const EmITreesNode xor_model_nodes[14] = {  
    { 1, 0.197349f, 1, 2 },  
    .....  
    { 0, 0.421164f, -2, -1 }  
};  
EmITrees xor_model = { 14, xor_model_nodes, ..... };  
  
static int32_t  
eml_trees_predict_tree(const EmITrees *f, int32_t tree_root,  
                      const float *features, int8_t features_length) {  
    int32_t n = tree_root;  
    while (n >= 0) {  
        const float value = f->nodes[node_idx].feature;  
        const float point = forest->nodes[node_idx].value;  
        const int16_t child = (value < point) ?  
            f->nodes[n].left : f->nodes[n].right;  
        ....  
    }  
    return leaf;  
}
```

Code generation (“inline”)

- + Supports int8/int16/int32 and float

```
static inline int32_t  
xor_model_tree_0(const float *features, int32_t  
features_length)  
{  
    if (features[1] < 0.197349f) {  
        if (features[0] < 0.466316f) {  
            return 0;  
        } else {  
            return 1;  
        }  
    } else {  
        if (features[1] < 0.256702f) {  
            if (features[0] < 0.464752f) {  
                return 0;  
            } else {  
                return 1;  
            }  
        } else {  
            ....  
        }  
    }  
}
```

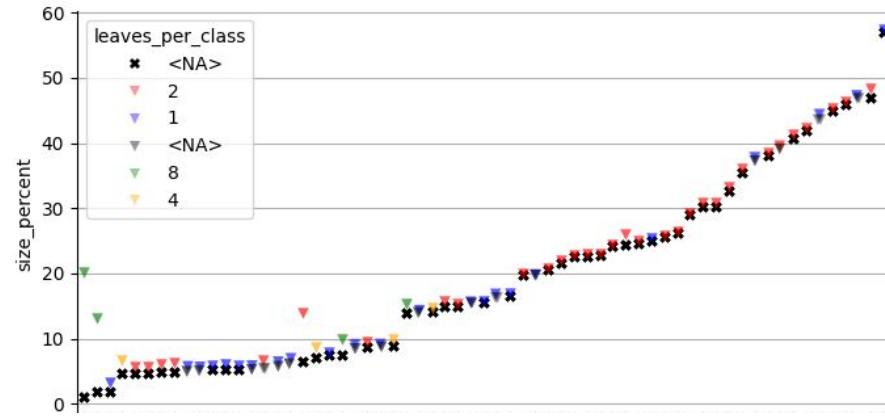
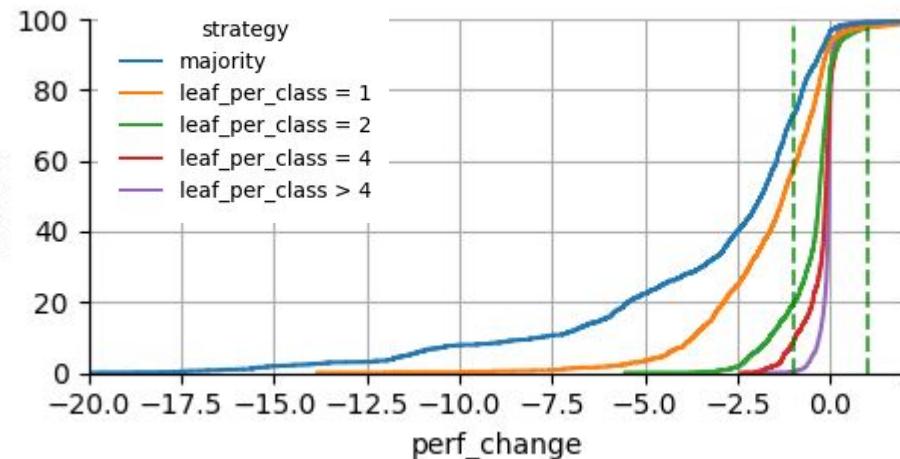
emlearn trees: K-means clustering for leaf de-duplication

Optimizing for: Small model size

Constraint: Under 1 percentage point drop in AUC ROC

Evaluating on 48 datasets from OpenML CC18 suite

Preliminary results
- paper to follow



Hard majority
sometimes bad:

Failed perf constraint on
75% of datasets

**Clustering with
N leaves per class**
Allows adaptable size/perf tradeoff

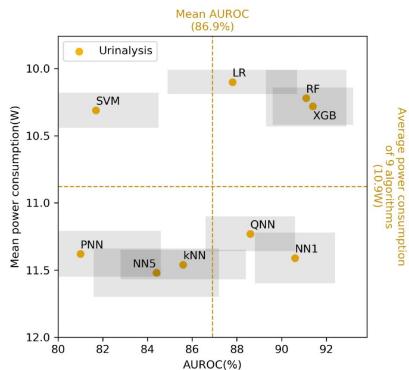
2-8 leaves per class sufficient
0% of datasets failed performance constraint
High model size compression 50-90%

Trees - relevant across the task complexity range

Low

Go-to for tabular data

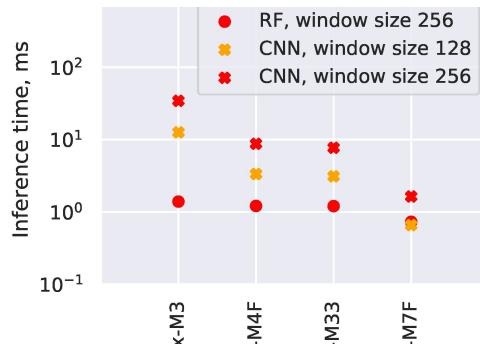
Very easy to get good performance, with minimal tweaking.



Medium

Alternative to deep learning

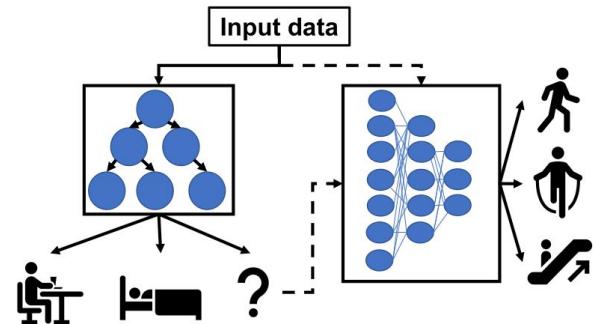
Elsts et.al [1] showed performance matching a deep-learning approach, at 10x to 100x better resource usage



High

Combine with neural network

Daghero et.al [2] showed up to 67.7% energy saving using two-stage approach



1. "Are Microcontrollers Ready for Deep Learning-Based Human Activity Recognition?"

Atis Elsts, Ryan McConville (2021) <https://www.mdpi.com/2079-9292/10/21/2640>

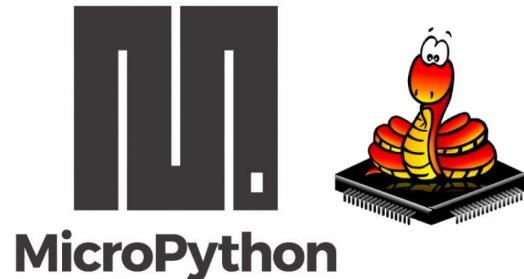
2. "Two-stage Human Activity Recognition on Microcontrollers with Decision Trees and CNNs".

Francesco Daghero, Daniele Jahier Pagliari, Massimo Poncino (2022) <https://arxiv.org/abs/2206.07652>

3. "Energy Efficiency of Inference Algorithms for Clinical Laboratory Data Sets: Green Artificial Intelligence Study"

Jia-Ruei Yu et. al (2022) <https://www.jmir.org/2022/1/e28036/>

More MicroPython



TinyML for MicroPython - comparisons

Project	Deployment	Models	Program size	Compute time
emlearn	Easy. Native mod .mpy	DT, RF, KNN, CNN	Good	Good
everywhereml	Easy. Pure Python .py	DT, RF, SVM, KNN,	High with large models	Poor
m2cgen	Easy. Pure Python .py	DT, RF, SVM, KNN, MLP	High with large models	Poor
OpenMV.tf	Hard. Custom Fork	CNN	High initial size	Good
ulab	Hard. User C module	<u>(build-your-own)</u> <u>Using ndarray</u> <u>primitives</u>	High initial size	Unknown (assume good)

MicroPython *is* Python - but not CPython

High degree of compatibility - but never 100%

Continuous job to keep up with CPython

Some differences inherent - from < 1 MB RAM and FLASH

Included libraries are minimal

micropython-lib has more extensive/featured

<https://github.com/micropython/micropython-lib>

Known incompatibilities

<https://docs.micropython.org/en/latest/genrst/index.html>

[#micropython-differences-from-cpython](#)

Not implemented (by CPython major release)

<https://github.com/micropython/micropython/issues/7919#issuecomment-1025221807>

No CFFI or C module compatibility!

But there is another C API

Garbage collected

One GC cycle *will take 1-10 ms* (typ)

Some control, limited (gc module)

TLDR:

- Will be very familiar to Python devs
- Small scripts will mostly work with minor mods.
- Larger programs/modules may need refactoring or rewrite to fit target

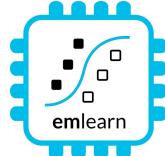
Hardware access - the **machine** module

<https://docs.micropython.org/en/latest/library/machine.html>

- class Pin – control I/O pins
- class Signal – control and sense external I/O devices
- class ADC – analog to digital conversion
- class ADCBlock – control ADC peripherals
- class PWM – pulse width modulation
- class UART – duplex serial communication bus
- class SPI – a Serial Peripheral Interface bus protocol (controller side)
- class I2C – a two-wire serial protocol
- class I2S – Inter-IC Sound bus protocol
- class RTC – real time clock
- class Timer – control hardware timers
- class WDT – watchdog timer
- class SD – secure digital memory card (cc3200 port only)
- class SDCard – secure digital memory card
- class USBDevice – USB Device driver

Hardware Abstraction Layer
for microcontroller peripherals

Same on all hardware/ports
* with exceptions



mip - package manager

Install from micropython-lib

mpremote install requests

Third party packages

mpremote install github:jonnor/micropython-zipfile

Can run directly on device *

```
import mip  
mip.install('requests')
```

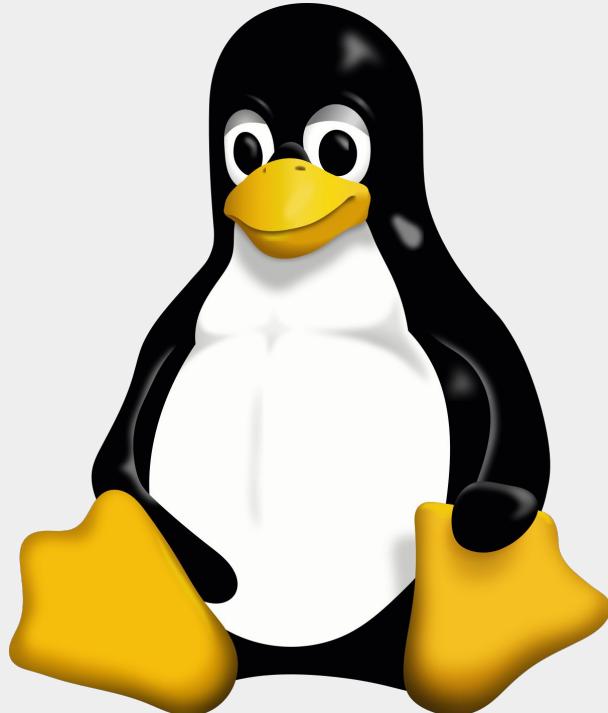
* Assuming device has Internet over WiFi/Ethernet

Install native C modules at runtime

```
mpremote mip install  
https://example.net/  
xtensawin_6.2*/emlearn_trees.mpy
```

* Specify architecture + MicroPython ABI version

MicroPython for Embedded Linux





Motivation: Memory efficiency

CPython is quite RAM hungry. Especially “standard” Python/PyData ecosystem

scikit-learn with CPython:

13 - 128 MB

```
from sklearn.ensemble import RandomForestClassifier  
estimator = RandomForestClassifier()  
# model not even trained yet!!!
```

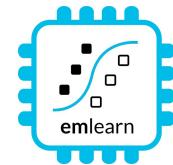
emlearn_trees with MicroPython:

0.1 - 6 MB

100 trees, 4000 nodes per tree (max_depth=12)

```
import emlearn_trees  
model = emlearn_trees.new(100, 4000, 100)
```

MicroPython attractive for Linux devices with < 512 MB RAM





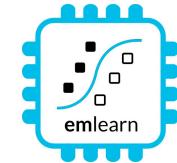
Unix MicroPython port: Limited hardware access

Not implemented:

- GPIO. machine.Pin / machine.PWM / machine.ADC
- Digital busses. machine.I2C / SPI / USART / USB
- Watchdog Timer. machine.WDT
- Power management. machine.lightsleep() /deepsleep()
- Audio input/microphone
- Camera access

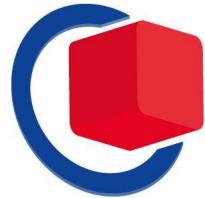
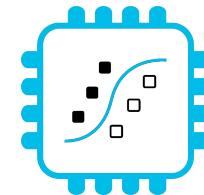
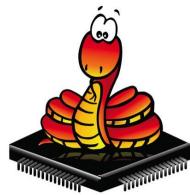
Need C or Python modules for this. Or call external programs.
Contributions welcome!

Already very useful for unit testing



Bridging the TinyML language gap

with MicroPython and emlearn



embeddedworld

Exhibition&Conference

Jon Nordby
jononor@gmail.com
March 2025, Nuremberg

<https://github.com/emlearn/emlearn>