


# Milliwatt sized machine learning

## on microcontrollers using emlearn

<https://github.com/emlearn/emlearn>



FOSDEM 2025, Brussels  
Low level AI hacking devroom  
Jon Nordby [jononor@gmail.com](mailto:jononor@gmail.com)





We utilise sound and vibration analysis to detect and warn you of upcoming errors in your technical infrastructure before they happen.

 **sound sensing**

## Condition Monitoring

Devices overview

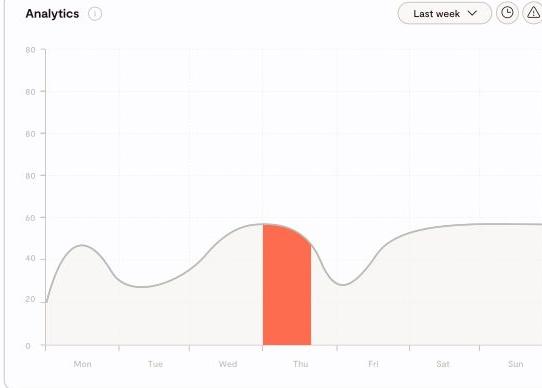
Search Filter by Organization Hotel Manana

**Tech\_Storage**

Location	Room	Status
Sandstuvein...	Island 04	<span style="color: green;">●</span>
Storgata 25...	TBJ 291	<span style="color: orange;">●</span>
Chr. Krohgs...	Tech 275	<span style="color: red;">●</span>
Møllergata 12...	Ohio 3	<span style="color: green;">●</span>
Ruseløkkveien...	HKM 261	<span style="color: green;">●</span>
Kristian IVs...	Freeway 273	<span style="color: orange;">●</span>
C. J. Hambro...	Oxaca2	<span style="color: green;">●</span>
Akersgata 65...	Storage_3	<span style="color: green;">●</span>

**Tech 275**

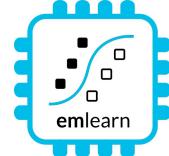
Analytics Last week



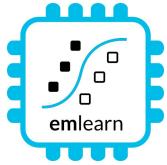
Mon Tue Wed Thu Fri Sat Sun

Trusted by Nordic market leaders

# Outline / agenda



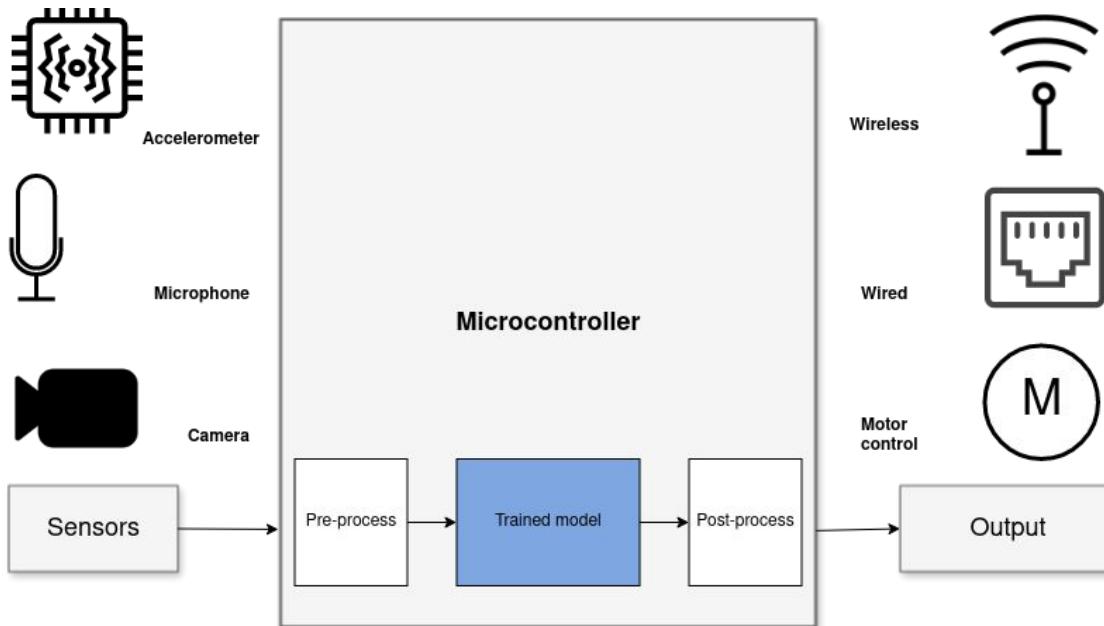
1. Why ML on microcontrollers (“**TinyML**”)
2. About the **emlearn project**
3. **Projects** made with emlearn
4. Quick how-to for **emlearn**



# Introduction

How is machine learning relevant  
for constrained embedded devices  
and microcontrollers?

# Sensor data analysis with Machine Learning



Benefits vs sending  
data to cloud

- Stand-alone
- Low latency
- Power efficiency
- Privacy
- Low cost

System diagram for “TinyML” solution

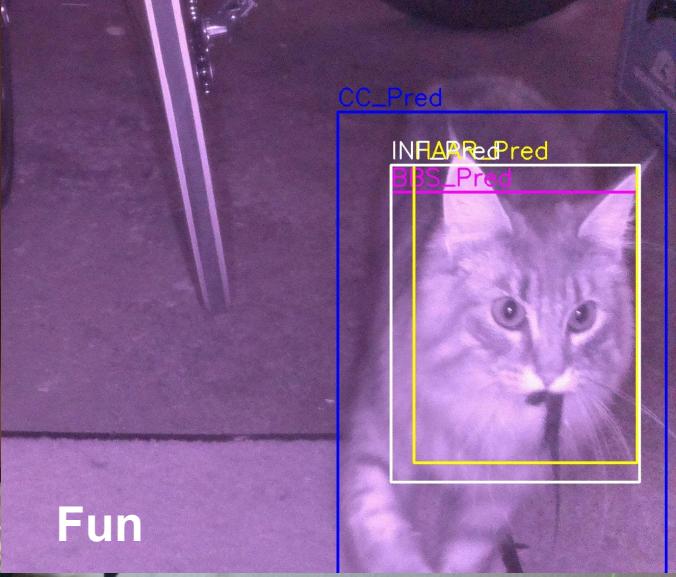
Hey Google...



Consumer Tech



Industrial



Fun



# Microcontroller - tiny programmable chip

Compute power: 1 / 1000x of a smartphone

- RAM: 0.10 - 1 000 kB
- Program space: 1.0 - 10 000 kB
- Compute 10 - 1000 DMIPS
- Price: 0.10 - 10 USD

Over 20 *billions* shipped per year

Increasingly accessible for hobbyists

2010: Arduino Uno

2016: ESP32

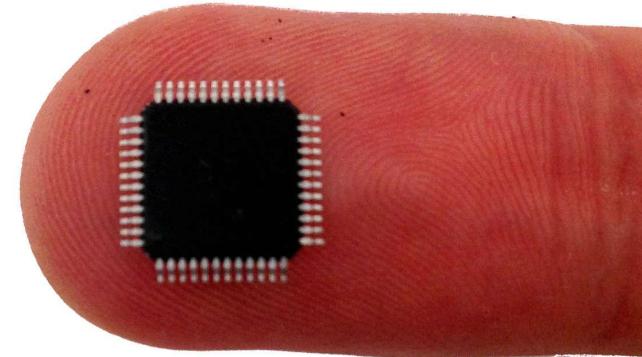

2021: RP2040

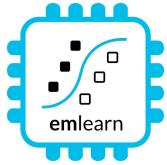
2024: RP2350

Arduino IDE C/C++

MicroPython

Efficiency is key !  
Memory, compute, power





# About emlearn

Started in 2018


Open-source (MIT)

## 1. Train on PC

pip install emlearn

### Convenient training

- Model creation in Python
- Use standard libraries
  - a. scikit-learn
  - b. Keras
- One-line to export to device
- Verification tools included




## 2. Deploy on device



### Embedded Friendly

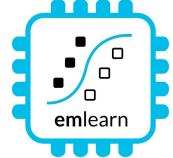
- Portable C99 code
- No dynamic allocations
- Header-only
- High test coverage
- Integer/fixed-point math \*
- Small. 2 kB+ FLASH




### Convenient & Efficient

- Portable MicroPython
- Single .mpy file
- No dynamic allocations
- High test coverage
- Integer/fixed-point math \*
- Small. 2 kB+ FLASH


# Supported tasks




## Anomaly Detection



## Classification



## Regression



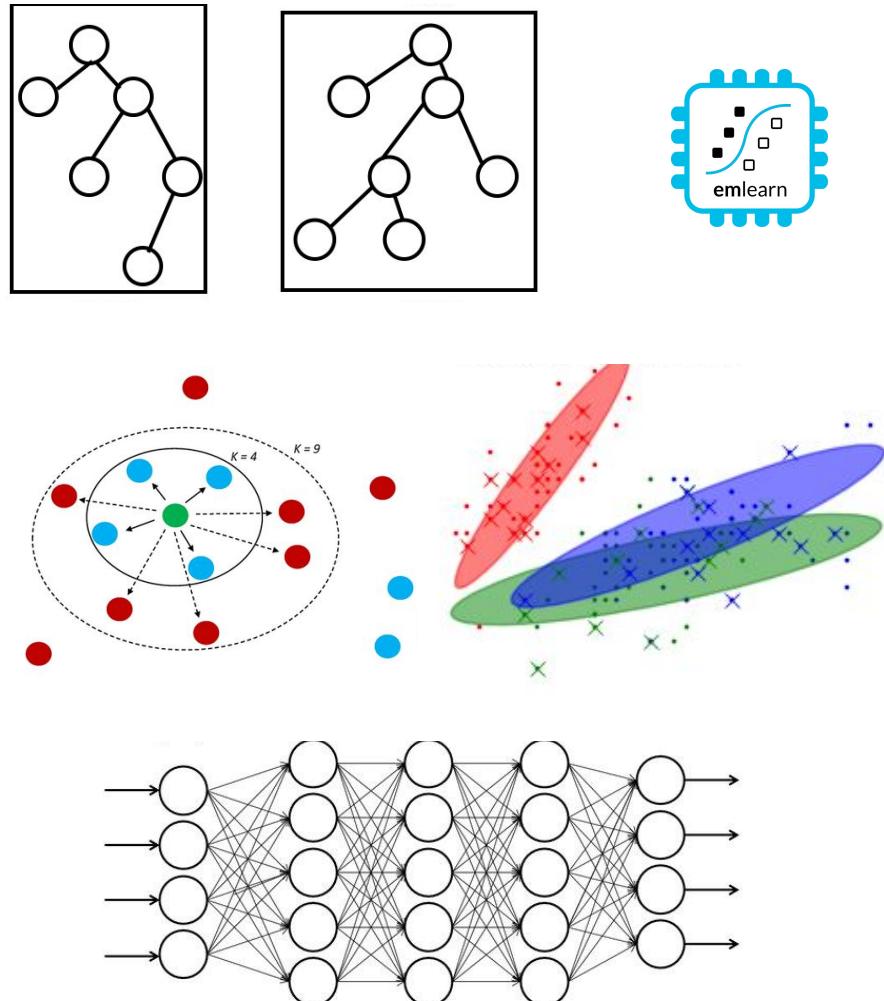
Supports the most common tasks  
for embedded  
& sensor data use cases.


- Classification
- Regression
- Anomaly Detection

# Supported models

Selection of simple & effective  
embedded-friendly models

- Decision Trees (DT)
- Random Forest (RF)
- K Nearest Neighbors (KNN)
- Gaussian Mixture Models (GMM)
- Multi-Layer-Perceptron (MLP)
- Convolutional Neural Network (CNN)





# Projects using emlearn

Many real-world uses  
across the globe

Referred in 40+ publications

# Health monitoring for cattle

Accelerometer data

used to detect abnormal behavior.

Can indicate **health issues** or other problems


Classified using a **Decision Tree**

*lying/walking/standing/grazing/ruminating*


Activity is transmitted using **LoRaWAN**

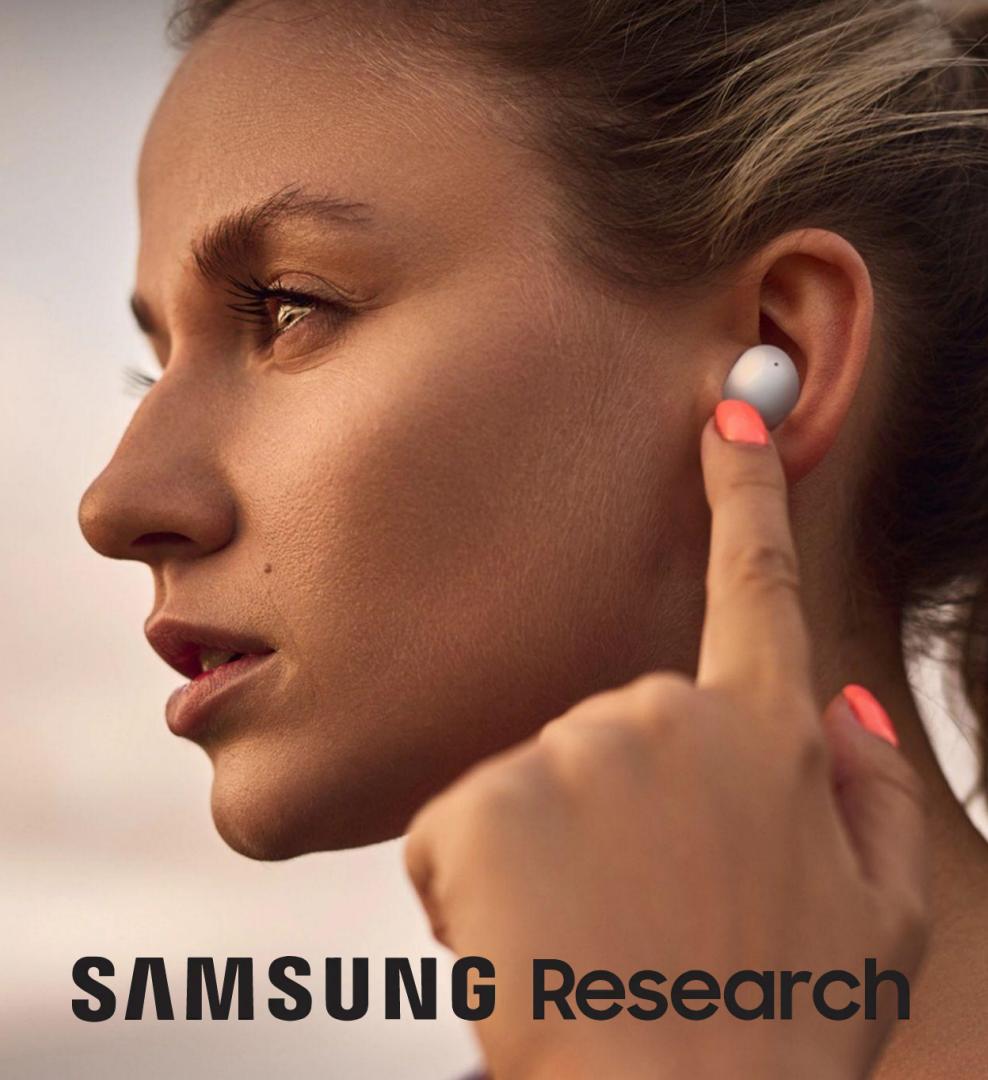
Running ML on-sensor: **under 1 mW**

**50 times lower power** than sending raw data



**Power Efficient Wireless Sensor Node through Edge Intelligence**  
Abhishek Damle (2022)





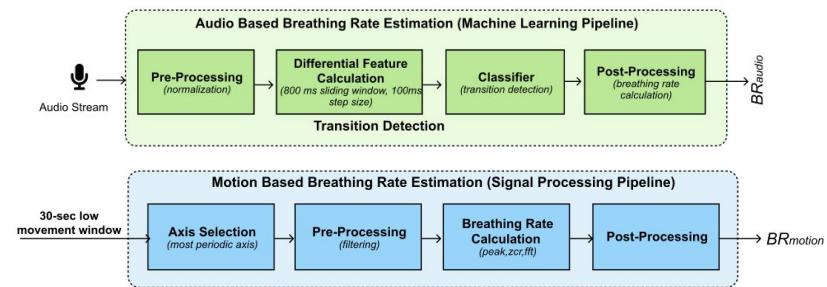
**SAMSUNG** Research

# Health monitoring earable

Breathing rate is critical to monitor for persons with **respiratory health problems**.

Monitoring **for every-day use**, not just in clinical context. Using earable to replace specialized devices.


Random Forest classifier for audio.



*Remote Breathing Rate Tracking in Stationary Position  
Using the Motion and Acoustic Sensors of Earables  
Tousif Ahmed et.al (2023)*

# Automatic toothbrush timer

<https://github.com/jonnor/toothbrush/>



- Stock hardware: M5StickC PLUS 2 by M5Stack
- ~500 lines of MicroPython code
- Collected and labeled data in ~1 hour

Data collection and training tools:

[https://github.com/emlearn/emlearn-micropython/tree/master/examples/har\\_trees](https://github.com/emlearn/emlearn-micropython/tree/master/examples/har_trees)



# 1 dollar TinyML

<https://hackaday.io/project/194511-1-dollar-tinyml>

## Research question:

Can we make a useful TinyML system  
<1 USD in total component cost (BOM)?




Preliminary answer: YES!

Either motion (accelerometer)  
or sound sensor (microphone)

Constraints: 8 kB RAM / 64 kB FLASH

## Current status:

Audio streaming works  
Revision 2 ordered





# emlearn for C

Quick walkthrough

# Train and export a model



## 1. Train using standard Python ML libraries.

```
from keras import ...      A) keras neural  
                           network  
  
model = Sequential([  
    Dense(16, input_dim=n_features, activation='relu'),  
    Dense(8, activation='relu'),  
    Dense(1, activation='sigmoid'),  
])  
model.compile(...)  
model.fit(X_train, Y_train, epochs=1, batch_size=10)
```

```
from sklearn.neural_network import MLPClassifier  
model = MLPClassifier(hidden_layer_sizes=(100,50,25))  
  
model.fit(X_train, Y_train)  
  
B) scikit-learn neural network
```

## 2. Use `emlearn.convert()` and `.save()`

```
import emlearn  
  
cmodel = emlearn.convert(model, method='inline')  
  
cmodel.save(file=mynet_model.h', name='mynet')
```



```
#include <eml_net.h>  
  
static const float mynet_layer_0_biases[8] = { -0.015587f, -0.005395f, -0.010957f, 0.015883f ....  
static const float mynet_layer_0_weights[24] = { -0.256981f, 0.041887f, 0.063659f, 0.011013f, ....  
static const float mynet_layer_1_biases[4] = { 0.001242f, 0.010440f, -0.005309f, -0.006540f };  
static const float mynet_layer_1_weights[32] = { -0.577215f, -0.674633f, -0.376140f, 0.646900f, ....  
static float mynet_buf1[8];  
static float mynet_buf2[8];  
static const EmlNetLayer mynet_layers[2] = {  
    { 8, 3, mynet_layer_0_weights, mynet_layer_0_biases, EmlNetActivationRelu },  
    { 4, 8, mynet_layer_1_weights, mynet_layer_1_biases, EmlNetActivationSoftmax }  
};  
static EmlNet mynet = { 2, mynet_layers, mynet_buf1, mynet_buf2, 8 };  
  
int32_t  
mynet_predict(const float *features, int32_t n_features)  
{  
    return eml_net_predict(&mynet, features, n_features);  
}  
.....
```

Example of generated code

# Using the C code



## 3. #include and call predict()

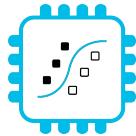
```
// Include the generated model code  
#include "mynet_model.h"  
  
// index for the class we are  
detecting  
#define MYNET_VOICE 1  
  
// Buffers for input data  
#define N_FEATURES 6  
float features[N_FEATURES];  
  
#define DATA_LENGTH 128  
int16_t  
sensor_data[DATA_LENGTH];
```

setup

```
// Get data and pre-process it  
read_microphone(sensor_data, DATA_LENGTH);  
preprocess_data(sensor_data, features);  
  
// Run the model  
out = mynet_predict(features, N_FEATURES);  
  
// Do something with results  
if (out == MYNET_VOICE) {  
    set_display("voice detected");  
} else {  
    set_display("");  
}
```

loop





**emlearn**  
-micropython

# emlearn for MicroPython

Python is the lingua franca  
for Machine Learning


How can we enable people to build  
TinyML solution also using Python?

# MicroPython - Introduction

Implements a subset of Python 3.x

For devices with 16 kB+ RAM

Supports 8+ microcontroller families



As compatible with CPython as possible, within constraints.

**No CFFI or C module compatibility!**

TLDR: Small .py scripts will mostly work with minor mods

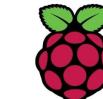
But not PyData stack

“mip” package manager

**Can load C modules at runtime!**

More info: <https://micropython.org>

RP2040



# MicroPython library for emlearn

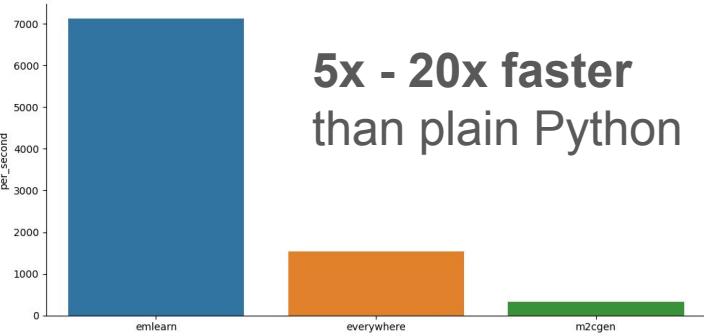


Combine the **convenience, familiarity and productivity of Python** - with the **speed of C**

Modules installable at runtime (.mpy)

emlearn\_iir  
emlearn\_trees  
emlearn\_fft  
emlearn\_cnn  
emlearn\_neighbors

Infinite Impulse Response filters  
Random Forest  
Fast Fourier Transform  
Convolutional Neural Networks  
K-nearest Neighbors



**5x - 20x faster**  
than plain Python

**Image classification+**  
**On-device learning**

<https://github.com/emlearn/emlearn-micropython/>

# Install & Export model



```
import emlearn
```

Export model on PC

```
estimator = train_model()  
converted = emlearn.convert(estimator)  
converted.save(name='gesture', format='csv', file='gesture_model.csv')
```

Copy model to device

```
mpremote fs cp gesture_model.csv :
```

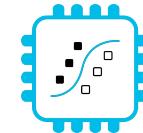
**Prebuilt binary**

No need to setup  
C toolchain and SDK

Copy library to device

```
mpremote mip install https://emlearn.github.io/....../xtensawin_6.2/emltrees.mpy
```

# Load model & run



emlearn  
-micropython

```
import emltrees
# Instantiate model. Capacity for trees, decision nodes, classes
model = emltrees.new(20, 200, 10)

# Load model "weights"
with open('gesture_model.csv') as f:
    emltrees.load_model(model, f)
```


Load model on device

```
# Read sensor data
sensor.fifo_read(samples)

# Preprocess features
processor.run(samples, features)


# Run model
out = model.predict(features)
```

Run inference




# Outro

# Summary



1. Machine Learning used in embedded systems  
**to automatically analyze sensor data**
2. Practical applications possible with just a  
**kilo-bytes RAM, milli-watts of power, couple dollars**
3. **emlearn** is an open-source project  
**to help deploy ML models to microcontrollers**  
- running **C or MicroPython**

# More resources



Official documentation

<https://emlearn-micropython.readthedocs.io>

<https://emlearn.readthedocs.io>

PyCon Berlin 2024: **Machine Learning on microcontrollers using MicroPython and emlearn**

<https://www.youtube.com/watch?v=S3GjLr0ZIE0>

TinyML EMEA 2024: **emlearn - scikit-learn for microcontrollers and embedded systems**

<https://www.youtube.com/watch?v=LyO5k1VMdOQ>


PyData ZA 2024: **Sensor data processing on microcontrollers with MicroPython**

<https://za.pycon.org/talks/31-sensor-data-processing-on-microcontrollers-with-micropython/>


# Milliwatt sized machine learning

## on microcontrollers using emlearn

<https://github.com/emlearn/emlearn>




FOSDEM 2025, Brussels  
Low level AI hacking devroom  
Jon Nordby [jononor@gmail.com](mailto:jononor@gmail.com)



# Bonus

# MicroPython - Installing



Download prebuilt firmware

<https://micropython.org/download/?port=esp32>

Flash firmware to device

pip install esptool


esptool.py --chip esp32 --port ... erase\_flash

esptool.py --chip esp32 --port ... write\_flash -z 0 micropython-v1.17.bin

Connect to device

pip install mpremote

mpremote repl



```
MicroPython v1.8.3-24-g095e43a on 2016-08-16; ESP module
Type "help()" for more information.
>>> print('Hello world!')
Hello world!
>>> █
```


IDE (optional): Thonny, VS Code, et.c.

# TinyML for MicroPython - comparisons

Project	Deployment	Models	Program size	Compute time
emlearn	Easy. Native mod .mpy	DT, RF, KNN, CNN	Good	Good
everywhereml	Easy. Pure Python .py	DT, RF, SVM, KNN,	High with large models	Poor
m2cgen	Easy. Pure Python .py	DT, RF, SVM, KNN, MLP	High with large models	Poor
OpenMV.tf	Hard. Custom Fork	CNN	High initial size	Good
ulab	Hard. User C module	<u>(build-your-own)</u> <u>Using ndarray</u> <u>primitives</u>	High initial size	Unknown (assume good)

# Task feasibility versus microcontroller type

	RAM (kB)	FLASH (kB)	CPU (CoreMark)		IMU 100 hz 5 s	Sound 16 kHz 1 s	Image 64x64 px 1 fps
<b>Arduino Uno</b> ATMega 328 @ 16 Mhz	2	32	4.0		✓	✗	✗
<b>Arduino Nano BLE 33</b> NRF52840 ARM Cortex M4F @ 64 Mhz	256	1 000	215		✓	✓	✗
<b>Arduino Nano ESP32</b> ESP32-S3 @ 240 Mhz	8 000	16 000	613 (per core)		✓	✓	✓
<b>Teensy 4.0</b> ARM Cortex M7 @ 600 Mhz	1 000	2 000	2313 (per core)		✓	✓	✓




# Tree-based ensembles

Random Forest / Decision Tree

The most popular model in emlearn

# emlearn trees: Leaf compression

scikit-learn decision trees store class probabilities in leaf-nodes. Size:  $n_{\text{leaves}} * n_{\text{classes}} * 4$  bytes  
Consequence: Leaves take a large amount of space



emlearn (since 2019) does **hard majority voting** instead  
And does **leaf-node de-duplication** across all trees

Size:  $n_{\text{classes}} * 1 * 1$  byte  
Saving 50-80% of space

Tradeoff: May lead to reduced predictive performance with some leaf values

Lossless	0.0	1.0	0.0	0.0	1
Not ideal	0.0	0.0	1.0	0.0	2
	0.0	0.4	0.6	0.0	
	0.15	0.20	0.30	0.15	

.....

# emlearn trees: Inference modes

## Data structure (“loadable”)

- + Can be loaded at runtime
- Only supports float

```
static const EmITreesNode xor_model_nodes[14] = {  
    { 1, 0.197349f, 1, 2 },  
    .....  
    { 0, 0.421164f, -2, -1 }  
};  
EmITrees xor_model = { 14, xor_model_nodes, ..... };  
  
static int32_t  
eml_trees_predict_tree(const EmITrees *f, int32_t tree_root,  
                      const float *features, int8_t features_length) {  
    int32_t n = tree_root;  
    while (n >= 0) {  
        const float value = f->nodes[node_idx].feature;  
        const float point = forest->nodes[node_idx].value;  
        const int16_t child = (value < point) ?  
            f->nodes[n].left : f->nodes[n].right;  
        ....  
    }  
    return leaf;  
}
```

## Code generation (“inline”)

- + Supports int8/int16/int32 and float

```
static inline int32_t  
xor_model_tree_0(const float *features, int32_t  
features_length)  
{  
    if (features[1] < 0.197349f) {  
        if (features[0] < 0.466316f) {  
            return 0;  
        } else {  
            return 1;  
        }  
    } else {  
        if (features[1] < 0.256702f) {  
            if (features[0] < 0.464752f) {  
                return 0;  
            } else {  
                return 1;  
            }  
        } else {  
            ....  
        }  
    }  
}
```



# emlearn trees: K-means clustering for leaf de-duplication

Optimizing for: Small model size

Constraint: Under 1 percentage point drop in AUC ROC

Evaluating on 48 datasets from OpenML CC18 suite

Preliminary results  
- paper to follow



**Hard majority**  
sometimes bad:

Failed perf constraint on  
75% of datasets

**Clustering with  
N leaves per class**  
Allows adaptable size/perf tradeoff


**2-8 leaves per class sufficient**  
0% of datasets failed performance constraint  
High model size compression 50-90%

# Trees - relevant across the task complexity range

## Low

Go-to for tabular data


Very easy to get good performance, with minimal tweaking.



## Medium

Alternative to deep learning


Elsts et.al [1] showed performance matching a deep-learning approach, at 10x to 100x better resource usage



## High

Combine with neural network

Daghero et.al [2] showed up to 67.7% energy saving using two-stage approach



1. "Are Microcontrollers Ready for Deep Learning-Based Human Activity Recognition?"

Atis Elsts, Ryan McConville (2021) <https://www.mdpi.com/2079-9292/10/21/2640>

2. "Two-stage Human Activity Recognition on Microcontrollers with Decision Trees and CNNs".

Francesco Daghero, Daniele Jahier Pagliari, Massimo Poncino (2022) <https://arxiv.org/abs/2206.07652>

3. "Energy Efficiency of Inference Algorithms for Clinical Laboratory Data Sets: Green Artificial Intelligence Study"

Jia-Ruei Yu et. al (2022) <https://www.jmir.org/2022/1/e28036/>