

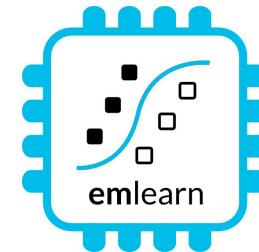
Microcontrollers + Machine Learning in 1-2-3

with MicroPython and emlearn

<https://github.com/emlearn/emlearn-micropython>



Jon Nordby jononor@gmail.com
PyData Global 2024





We utilise sound and vibration analysis to detect and warn you of upcoming errors in your technical infrastructure before they happen.

 **sound sensing**

Condition Monitoring

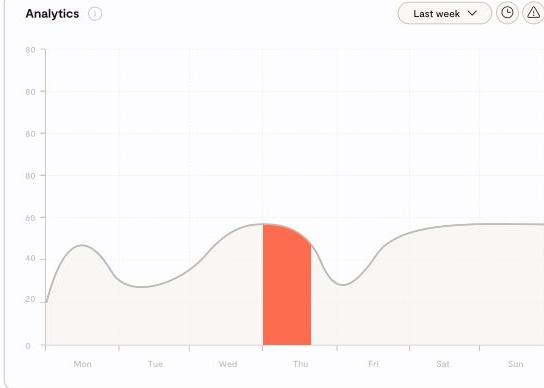
Devices overview Search Filter by Organization Hotel Manana

Tech_Storage

Location	Room	Status
Sandstuvein...	Island 04	Green
Storgata 25...	TBJ 291	Yellow
Chr. Krohgs...	Tech 275	Red
Møllergata 12...	Ohio 3	Green
Ruseløkkveien...	HKM 261	Green
Kristian IVs...	Freeway 273	Yellow
C. J. Hambro...	Oxaca2	Green
Akersgata 65...	Storage_3	Green

Tech 275

Analytics Last week More Less



Mon Tue Wed Thu Fri Sat Sun

Trusted by Nordic market leaders

Goal

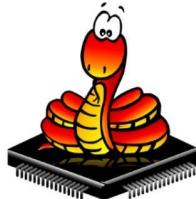
Purpose of this presentation

You, as a **Python developer**,
can **build an IoT sensor**
on a **microcontroller**
using **MicroPython**

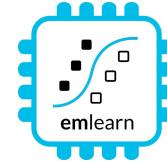
+ Including
on-edge Machine Learning
with **emlearn-micropython**

Python on Microcontroller

Jumping right into it



MicroPython



What is a microcontroller?

Modern microcontroller:
A complete programmable System-on-Chip

Example: ESP32-S3FH4R2

32 bit CPU, 240 Mhz

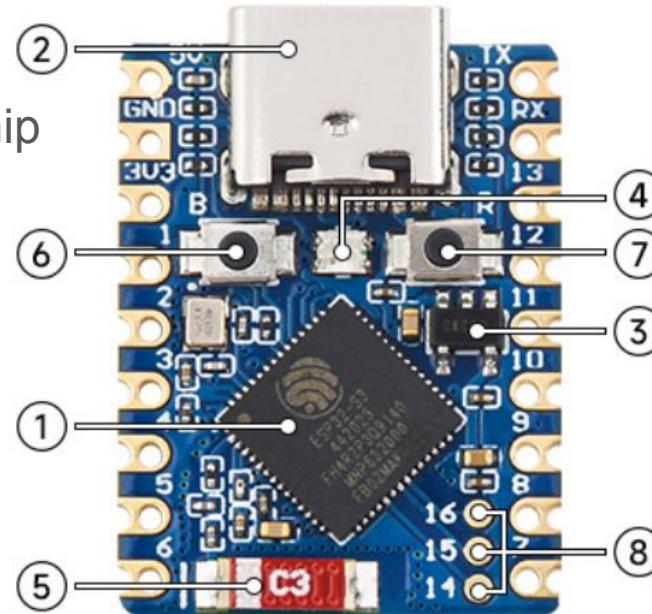
Floating Point Unit

2 MB RAM

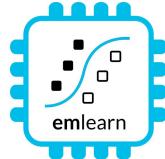
4 MB FLASH

WiFi

Bluetooth Low Energy
USB-C



Espressif ESP32-S3FH4R2 chip:	2.5	USD
Waveshare ESP32-S3-Tiny board:	6	USD



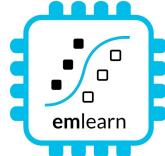
Hardware tinkering optional

- complete devices available



Complete ESP32 based device with sensors etc.:

20 - 50 USD



Installing MicroPython

Download prebuilt firmware

<https://micropython.org/download/?port=esp32>

Flash firmware to device

pip install esptool

esptool.py --chip esp32 --port ... erase_flash

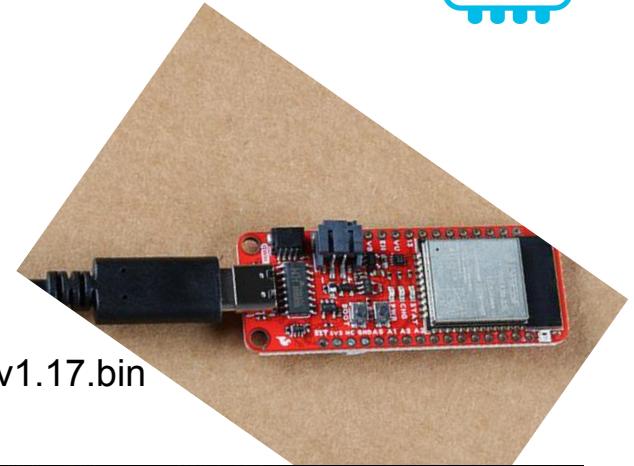
esptool.py --chip esp32 --port ... write_flash -z 0 micropython-v1.17.bin

Connect to device

pip install mpremote

mpremote repl

```
MicroPython v1.8.3-24-g095e43a on 2016-08-16; ESP module
Type "help()" for more information.
>>> print('Hello world!')
Hello world!
>>> █
```



IDE (optional): Viper IDE, Thonny, et.c.

Temperature sensor - hardware



I2C
INTERFACE

Example
MPU6886 accelerometer
built-in temperature sensor

Temperature Sensor
DS1820
Water-proof

Room Monitoring
+ CO2
+ humidity
+ temperature

Temperature sensor - code

Using <https://viper-ide.org/>

Zero-install. Connect to device via USB

1. Read the sensor in a loop
2. Send data using MQTT
3. Wait until next measurement

The same approach can be used for sensing other slow changing phenomena

```
test.py • X
1  from mqtt_as import MQTTClient, config
2  import asyncio
3  from mpu6886 import MPU6886
4  from machine import I2C
5
6  # Local configuration
7  config['ssid'] = 'FIXME' # Optional on ESP8266
8  config['wifi_pw'] = 'FIXME'
9  config['server'] = 'test.mosquitto.org'
10
11 mpu = MPU6886(I2C(0, sda=21, scl=22, freq=100000))
12
13 v async def main(client):
14     print('main-start')
15     await client.connect()
16     print('connected')
17
18 v     while True:
19         t = mpu.temperature
20         print('publish-data', t)
21         await client.publish('pydataglobal2024/send', f'{t:.2f}', qos = 0)
22         await asyncio.sleep(30)
23
24 MQTTClient.DEBUG = True # Optional: print diagnostic messages
25 client = MQTTClient(config)
26 v try:
27     asyncio.run(main(client))
28 v finally:
29     client.close()
```

Using [peterhinch/micropython-mqtt](https://github.com/peterhinch/micropython-mqtt) and [jonnor/micropython-mpu6886](https://github.com/jonnor/micropython-mpu6886)

Installing packages

MicroPython has a package manager “**mip**”

Directly on device!

ViperIDE supports install via UI

Custom packages via URL

The screenshot shows the ViperIDE interface. On the left, a sidebar titled "Tools" lists various options: "device connection" (WebREPL over internet, WebREPL in local network, Bluetooth REPL, P2P Bridge), "package manager" (Install package via link, Python), "Python" (Prettyfy current file, Minify current file). A central modal dialog box is open, prompting the user to "Enter package name or URL:" with a text input field containing "https://github.com/jonnor/micropython-mpu6886/" and two buttons: "Cancel" and "OK". Below the modal, the main workspace displays a block of MicroPython code:

```
7 config['ssid'] = 'FIXME' # Optional on ESP8266
8 config['wifi_pw'] = 'FIXME'
9 config['server'] = 'test.mosquitto.org'
10
11 mpu = MPU6886(I2C(0, sda=21, scl=22, freq=100000))
12
13 async def main(client):
14     print('main-start')
15     await client.connect()
16     print('connected')
17
18     while True:
19         t = mpu.temperature
20         print('publish-data', t)
21         await client.publish('pydataglobal2024/send', f'{t:.2f}', qos = 0)
22         await asyncio.sleep(30)
23
24 MQTTClient.DEBUG = True # Optional: print diagnostic messages
25 client = MQTTClient(config)
26 try:
27     asyncio.run(main(client))
28 finally:
29     client.close()
```

At the bottom of the screen, there is a terminal window showing the MicroPython REPL output:

```
Type "help()" for more information.
>>>
MicroPython v1.23.0-3.g2ceccbb531.dirty on 2024-09-25; Generic ESP32 module with SPIRAM with ESP32
Type "help()" for more information.
>>> >ESP32 module with SPIRAM with ESP32
Type "help()" for more information.
>>> summary-queue-overflow
summary-queue-overflow
```

Sensor nodes with MicroPython

What can be done on an ESP32
microcontroller with MicroPython
and how to make it work

Temperature logger

Temperature
0.1 Hz



Feasibility

Pure Python

Activity tracker

Accelerometer
100 Hz



Noise monitor

Microphone
16000 Hz



Python with C modules

Image Classifier

Camera
27000 bytes/s

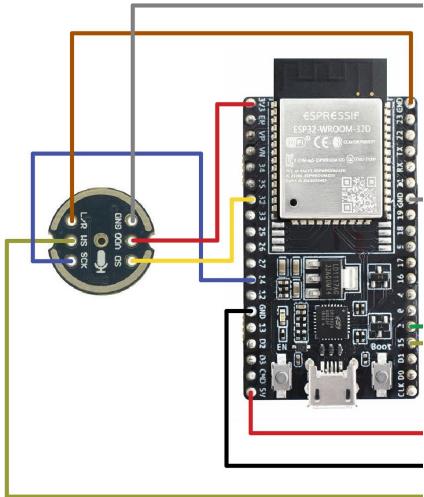


if is_MyCat(img):
open_door()

Sound
sensor

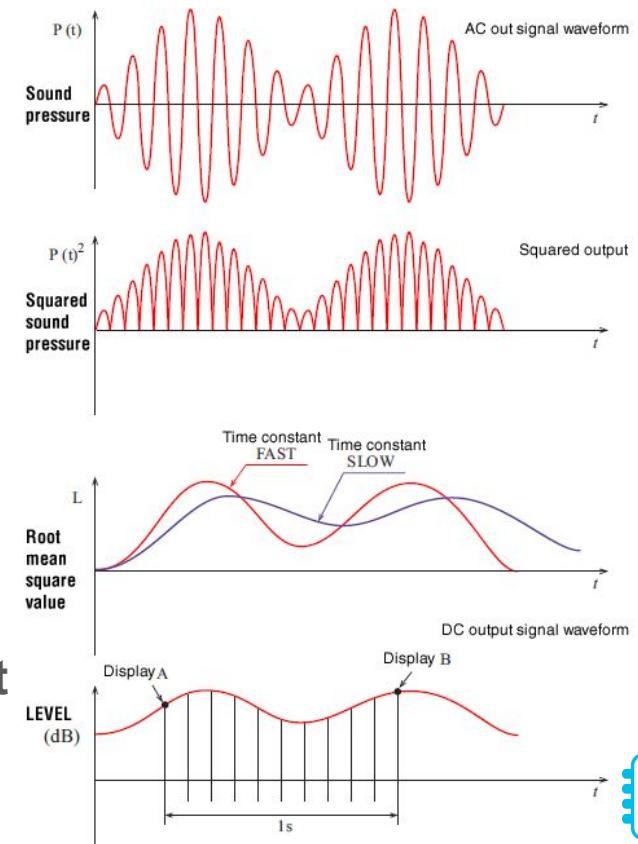
Sound sensor

I2S digital microphone (Example wiring)



Complete hardware unit
LilyGo TTGO
T-Camera S3

Soundlevel calculation Processing steps



Sound sensor - Install it

https://github.com/emlearn/emlearn-micropython/tree/master/examples/soundlevel_iir

Running on device (ViperIDE)

The fastest and easiest way to install on your device is to use Viper IDE. This will install the library and the example code automatically.

 Run in ViperIDE

In Viper IDE, you can select which example file to run (described below), and hit the Play button to run it.

Rendering text/widgets to screen

<https://github.com/peterhinch/micropython-nano-gui>

```
# I2S audio input
from machine import I2S
audio_in = I2S(0, sck=Pin(26), ws=Pin(32), sd=Pin(33),
               mode=I2S.RX, bits=16, format=I2S.MONO, rate=16000,
               )

# allocate sample arrays
chunk_samples = int(AUDIO_SAMPLERATE * 0.125)
mic_samples = array.array('h', (0 for _ in range(chunk_samples))) # int16
# memoryview used to reduce heap allocation in while loop
mic_samples_mv = memoryview(mic_samples)
# global to share state between callback and main
soundlevel_db = 0.0

meter = SoundlevelMeter(buffer_size=chunk_samples, samplerate=16000)

def audio_ready_callback(arg):
    # compute soundlevel
    global soundlevel_db
    soundlevel_db = meter.process(mic_samples)
    # re-trigger audio callback
    _ = audio_in.readinto(mic_samples_mv)

def main():
    # Use Non-Blocking I/O with callback
    audio_in.irq(audio_ready_callback)
    # Trigger first audio readout
    audio_in.readinto(mic_samples_mv)

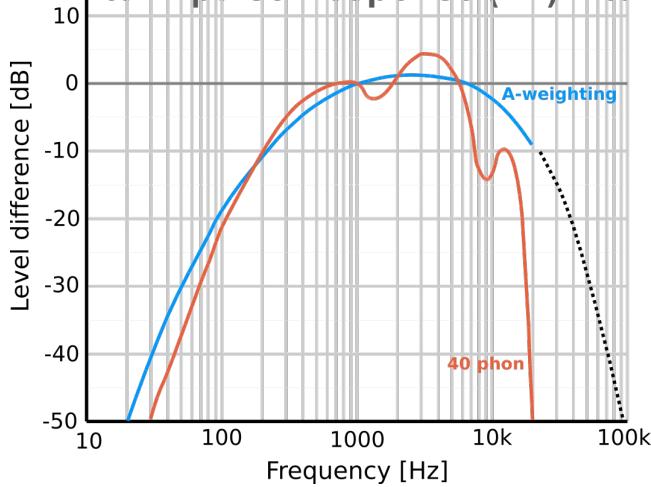
    while True:
        render_display(db=soundlevel_db)
        time.sleep_ms(200)

if __name__ == '__main__':
    main()
```

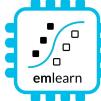
Sound sensor - IIR filter using emlearn_iir

Standard sound level measurements are **A-weighted**. To approximate human hearing.

Typically, implemented using
Infinite Impulse Response (IIR) filters.



Using `machine.I2S`
16 kHz samplerate



C modules

Written in C.

Defines a Python module with API.
functions/classes etc.

Can be implemented by users,
libraries or be part of MicroPython
core.

Can be portable or specific to one
hardware/platform

```
// Include the header file to get access to the MicroPython API
#include "py/dynruntime.h"

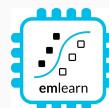
// Helper C function to compute factorial
static mp_int_t factorial_helper(mp_int_t x) {
    if (x == 0) {
        return 1;
    }
    return x * factorial_helper(x - 1);
}

// DEFINE FUNCTION. Callable from Python
static mp_obj_t factorial(mp_obj_t x_obj) {
    mp_int_t x = mp_obj_get_int(x_obj);
    mp_int_t result = factorial_helper(x);
    return mp_obj_new_int(result);
}
static MP_DEFINE_CONST_FUN_OBJ_1(factorial_obj, factorial);

// MODULE ENTRY
mp_obj_t mpy_init(mp_obj_fun_bc_t *self, size_t n_args, size_t n_kw, mp_obj_t *args) {
    // Must be first, it sets up the globals dict and other things
    MP_DYNRUNTIME_INIT_ENTRY

    // Register function in the module's namespace
    mp_store_global(MP_QSTR_factorial, MP_OBJ_FROM_PTR(&factorial_obj));

    // This must be last, it restores the globals dict
    MP_DYNRUNTIME_INIT_EXIT
}
```

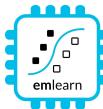


Native module (.mpy) VS External C module

	Native module	External C module
Installable at runtime	Yes, as .mpy file	No. Must be included in firmware image
Requires SDK/toolchain	No (only to build)	Yes
Code executes from	RAM	FLASH
Limitations	No libc / libm linked * No static BSS *	None
Maturity	Low *	Excellent
Documentation	https://docs.micropython.org/en/latest/develop/natmod.html	https://docs.micropython.org/en/latest/develop/cmodules.html

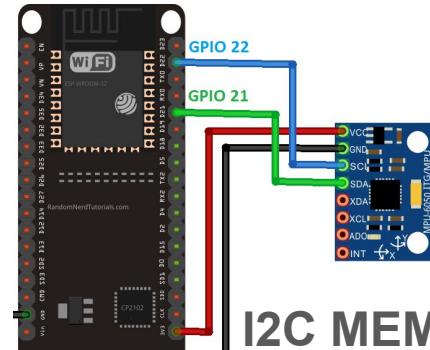
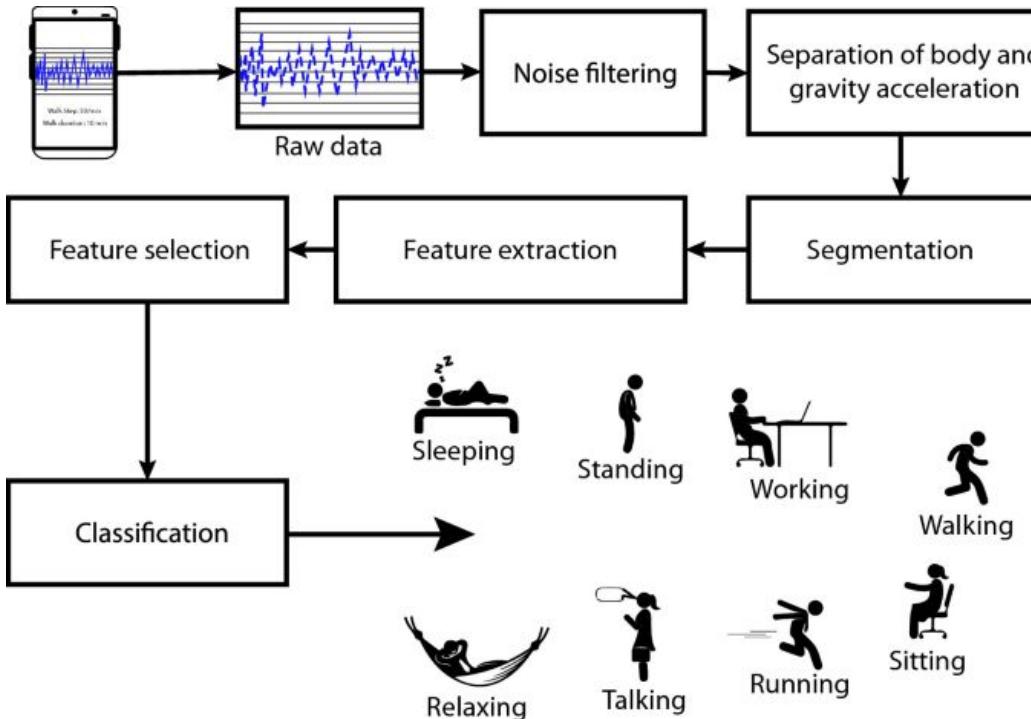
* Improved greatly in upcoming MicroPython (1.25+).

Contributions by Volodymyr Shymanskyy, Alessandro Gatti, Damien George, and others



Activity tracker

Activity tracker - concept



I2C MEMS IMU
accelerometer/gyro
(Example wiring)

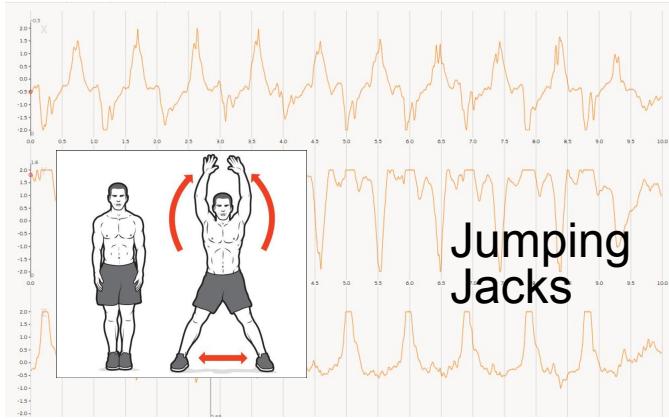


Complete
hardware unit
LilyGo
T-Watch S3

Recording a dataset

har_record.py

https://github.com/emlearn/emlearn-micropython/tree/master/examples/har_trees

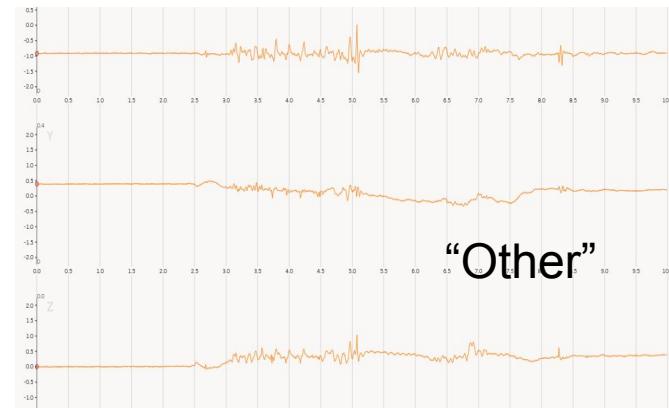


3 separate series,
1 minute per activity

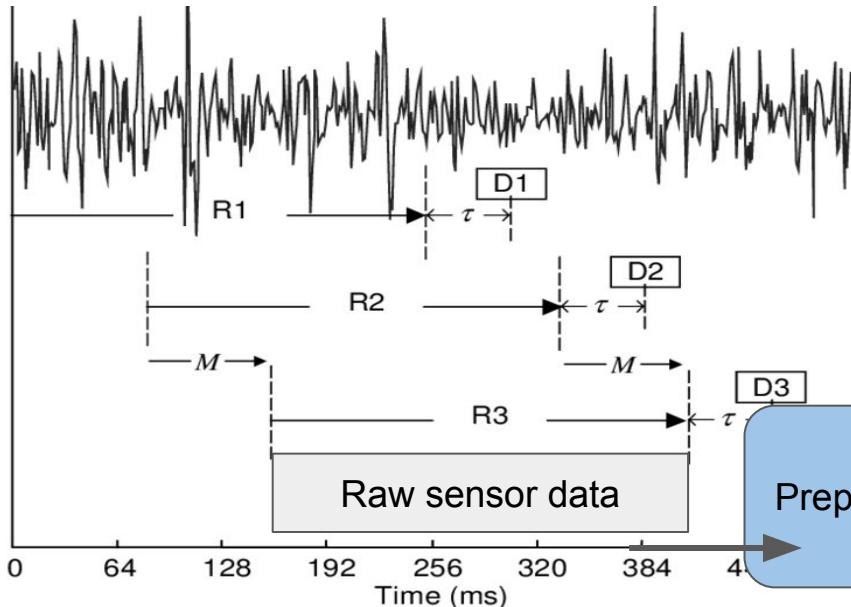
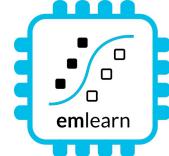
Saves 10 sec .npy files

Transfer over USB

Cleaned labels in
Label Studio



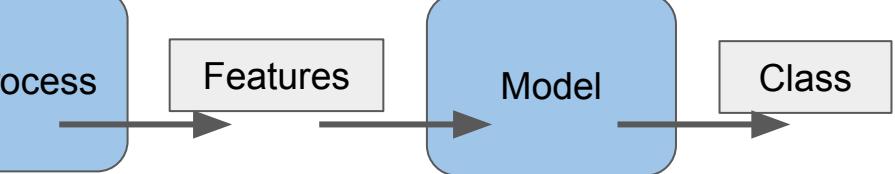
ML on streams: Continuous classification



The sensor data stream
is sliced into overlapping windows.
Each window processed independently

Exercise activity detection:

- 4 second window, every 1 second
- 100 Hz samplerate
- Processing time 200 ms



[42, 4002, ..., 329]

“Jumping Jacks”

Implementing an IMU/accelerometer/gyro driver? Use the FIFO!
<https://github.com/orgs/micropython/discussions/15512>

Activity Tracker - Feature Extraction

Statistical summarizations are useful time-series features, sufficient for basic Human Activity Recognition.

! Preprocessing *must be compatible*
between training on host PC (CPython) and device (MicroPython)

Solution: Write preprocessor for MicroPython, re-use in Python

```
subprocess('micropython preprocess.py data.npy features.npy')
```

Alternative: (when using common MicroPython/CPython subset)

```
import mypreprocessor.py
```

Using micropython-npyfile to read/write Numpy .npy files
<https://github.com/jonnor/micropython-npyfile/>

```
l = sorted(list(v))
l2 = [x*x for x in l]
sm = sum(l)
sqsum = sum(l2)
avg = sum(l) / len(l)

median = l[MEDIAN]
q25 = l[Q1]
q75 = l[Q3]
iqr = (l[Q3] - l[Q1])

energy = ((sqsum / len(l2)) ** 0.5)
std = ((sqsum - avg * avg) ** 0.5)
```

[https://github.com/emlearn/emlearn-micropython
/blob/master/examples/har_trees/timebased.py](https://github.com/emlearn/emlearn-micropython/blob/master/examples/har_trees/timebased.py)

Time-based features extraction
Are Microcontrollers Ready for Deep
Learning-Based Human Activity Recognition?
Atis Elsts, and Ryan McConville
<https://www.mdpi.com/2079-9292/10/21/2640>

Training model on dataset

Using a scikit-learn based pipeline.

```
# Setup subject-based cross validation
splitter = GroupShuffleSplit(n_splits=n_splits, test_size=0.25,
                             random_state=random_state)

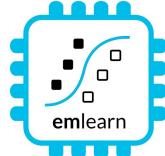
# Random Forest classifier
clf = RandomForestClassifier(random_state = random_state,
                            n_jobs=1, class_weight = "balanced")

# Hyper-parameter search
search = GridSearchCV(clf, param_grid=hyperparameters,
                      scoring=metric, refit=metric, cv=splitter)
search.fit(X, Y, groups=groups)
```

```
(venv) [jon@jon-thinkpad har_trees]$ MIN_SAMPLES_LEAF=150,200,400 python har_train.py
-dataset har_exercise_1 --window-length 400 --window-hop 10
2024-12-04 12:54:52 [info] data-loaded dataset=har_exercise_1
duration=0.016095876693725586 samples=32000
2024-12-04 12:54:56 [info] feature-extraction-done dataset=har_exercise_1
duration=4.534412145614624 labeled_instances=1952 total_instances=1952
Class distribution
activity
jumpingjack    549
lunge           488
other            488
squat            427
Name: count, dtype: int64
Model written to ./har_exercise_1_trees.csv
Testdata written to ./har_exercise_1.testdata.npz
Results
   n_estimators  min_samples_leaf  mean_train_f1_micro  mean_test_f1_micro
0             10              150          0.996311        0.962705
1             10              200          0.995628        0.956557
2             10              400          0.986202        0.920902
```

har_train.py

```
import emlearn
converted = emlearn.convert(clf)
converted.save(name='gesture', format='csv', file='model.csv')
```



Deploying trained model

Copy model to device

```
mpremote fs cp gesture_model.csv :
```

Load and run on device

```
import emltrees
model = emltrees.new(10, 1000, 10)
with open('eml_digits.csv', 'r') as f:
    emltrees.load_model(model, f)

features = array.array('h', ...)
out = model.predict(features)
```

Performance comparison

10 trees, max 100 leaf nodes, “digits” dataset

	Inference time	Program space
m2cgen	60.1 ms	179 kB
everywhereml	17.7 ms	154 kB
emlearn	1.3 ms	15 kB

emlearn is **10x faster and 10x more space efficient** compared to generating Python code

Summary

1. Modern microcontrollers are very accessible

Runs (Micro)Python!

ESP32 recommended

ViperIDE easy start

Conclusions

2. MicroPython productive environment for sensor devices

Python familiarity and ease-of-use

mip package manager

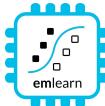
Good connectivity

3. Can implement advanced processing of sensor data

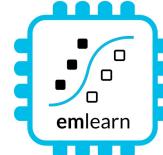
Accelerometer, audio, image, radar,

C modules a killer feature

emlearn-micropython: modules for DSP and Machine Learning



More resources



emlearn-micropython

`emlearn_iir`

`emlearn_trees`

`emlearn_fft`

`emlearn_tinymaix_cnn`

`emlearn_neighbors`

Infinite Impulse Response filters

Random Forest

Fast Fourier Transform

Convolutional Neural Networks

K-nearest Neighbors

**Image classification+
On-device learning**

Ulab: Numpy implementation for MicroPython

<https://github.com/v923z/micropython-ulab>

OpenMV: Computer/Machine Vision for MicroPython

<https://openmv.io/>

PyCon Berlin 2024: Machine Learning on microcontrollers using MicroPython and emlearn

<https://www.youtube.com/watch?v=S3GjLr0ZIE0>

PyData ZA 2024: Sensor data processing on microcontrollers with MicroPython

<https://za.pycon.org/talks/31-sensor-data-processing-on-microcontrollers-with-micropython/>

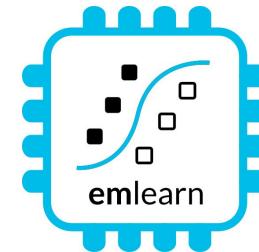
Microcontrollers + Machine Learning in 1-2-3

with MicroPython and emlearn

<https://github.com/emlearn/emlearn-micropython>

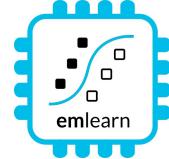


Jon Nordby jononor@gmail.com
PyData Global 2024



Bonus

TinyML for MicroPython - comparisons



Project	Deployment	Models	Size	Compute time
emlearn -micropython	Easy. Native mod .mpy	DT, RF, KNN, CNN	Good	Good
everywhereml	Easy. Pure Python .py	DT, RF, SVM, KNN,	High with large models	Poor
m2cgen	Easy. Pure Python .py	DT, RF, SVM, KNN, MLP	High with large models	Poor
OpenMV.tf	Hard. Custom Fork	CNN	High initial size	Good
ulab	Hard. User C module	(build-your-own) Using ndarray primitives	High initial size	Unknown (assume good)

Microcontroller - tiny programmable chip

Compute power: 1 / 1000x of a smartphone

- RAM: 0.10 - 1 000 kB
- Program space: 1.0 - 10 000 kB
- Compute 10 - 1 000 DMIPS
- Price: 0.10 - 10 USD
- Energy: 1 - 1 000 milliWatt

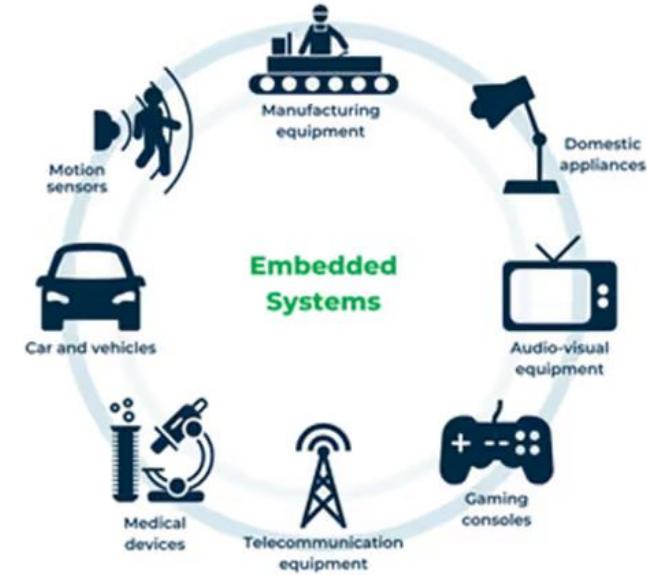
Over 20 *billions* shipped per year!

Increasingly accessible for hobbyists

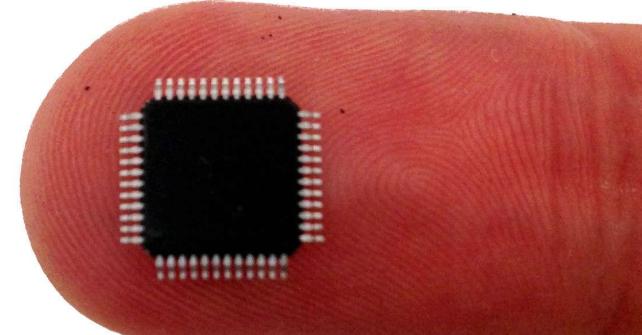
2010: Arduino Uno

2014: MicroPython

2019: MicroPython 1.10 - ESP32 PSRAM



Efficiency is key !
Memory, compute, power



Inline Assembly

MicroPython can expose Assembler opcodes as Python statements.

Allows to write a function in Assembler *inline in the Python program*

Can compile and execute on device

Supported on ARM Cortex M chips
Not supported (yet) on ESP32

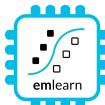
For the most hardcore hackers!

Official Documentation:

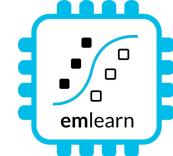
https://docs.micropython.org/en/latest/reference/asm_thumb2_index.html

```
@micropython.asm_thumb
def fir(r0, r1, r2):
    mov(r3, r8)                      # For Pico: can't push({r8}). r0-r7 only.
    push({r3})
    ldr(r7, [r0, 0])                 # Array length
    mov(r6, r7)
    mov(r3, r0)
    add(r3, 12)                     # r3 points to ring buffer start
    sub(r7, 1)
    add(r7, r7, r7)
    add(r7, r7, r7)
    add(r5, r7, r3)
    ldr(r4, [r0, 8])
    cmp(r4, 0)
    bne(INITIALISED)
    mov(r4, r3)
    label(INITIALISED)
    str(r2, [r4, 0])
    add(r4, 4)
    cmp(r4, r5)
    ble(BUFOK)
    mov(r4, r3)
    label(BUFOK)
    str(r4, [r0, 8])                # Save the insertion point for next call
    # *** Filter ***
    ldr(r0, [r0, 4])                # Bits to shift
```

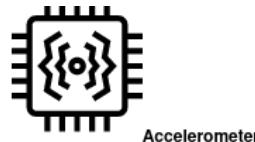
Example: FIR filter implementation (cut out)
<https://github.com/peterhinch/micropython-filters/blob/master/fir.py>



Sensor node systems



Including on-sensor data processing
with Digital Signal Processing (DSP) and Machine Learning (ML)



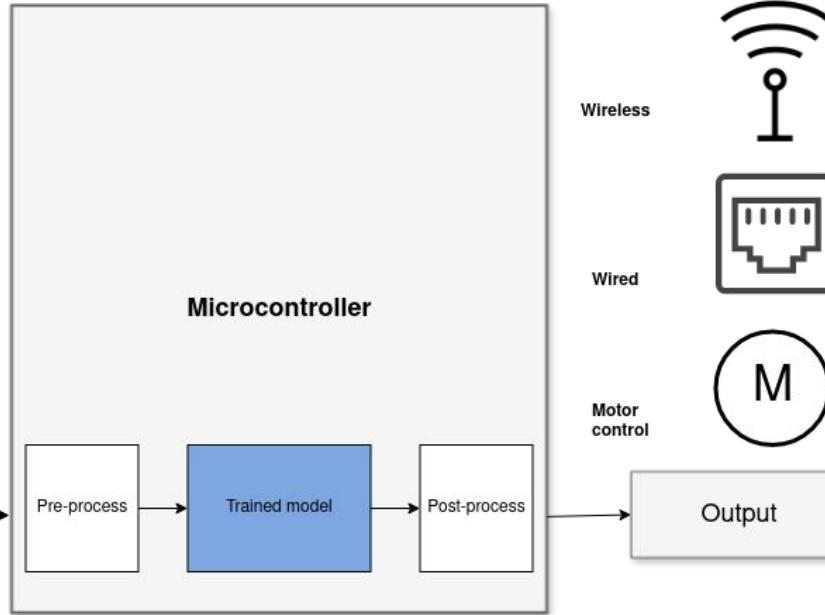
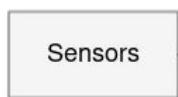
Accelerometer



Microphone



Camera



“Internet of Things”



“Wearable devices”



**Robotics,
Industrial automation**

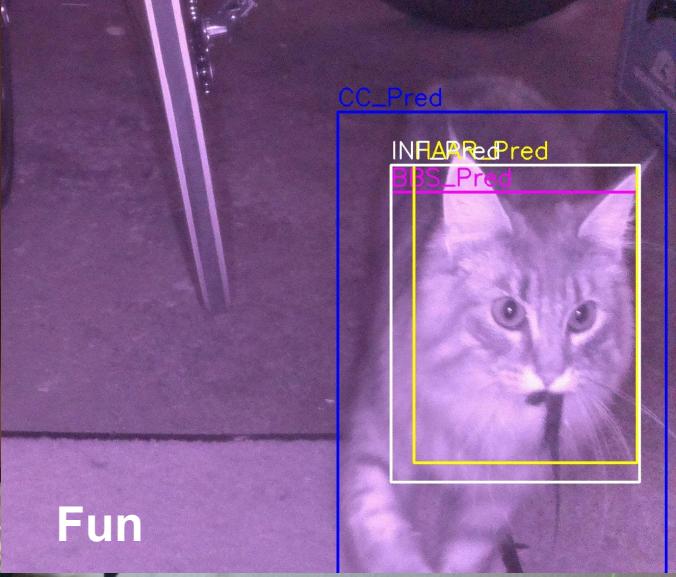
Hey Google...



Consumer Tech



Industrial



Fun



Cat Detector - Data Acquisition

Classify *low-res* images, at 1 FPS+

Using mp_camera by cnadler86
[https://github.com/
cnadler86/micropython-camera-API](https://github.com/cnadler86/micropython-camera-API)

96x96 px grayscale.
Takes ~100 ms - OK



Cat Detector - Classification with CNN

emlearn_tinyai_x_cnn

Go-to solution for image classification:
Convolutional Neural Network (CNN)

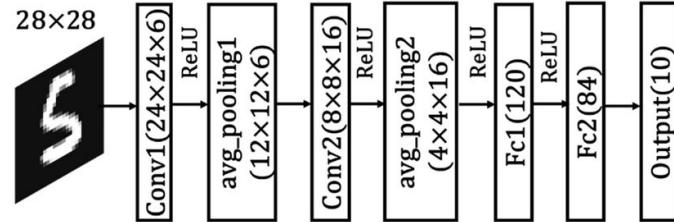
32x32 px input. 3 layers. Under 100 ms - OK

Preprocessing (untested)

Downscaling 96 px -> 32 px

Image brightness normalization

See also OpenMV - MicroPython for Computer Vision- <https://docs.openmv.io/index.html>



https://github.com/emlearn/emlearn-micropython/tree/master/examples/mnist_cnn

```
import tinyai_x_cnn # from emlearn-micropython

with open('cat_classifier.tmdl', 'rb') as f:
    model_data = array.array('B', f.read())
    model = tinyai_x_cnn.new(model_data)

while True:

    raw = read_camera()
    img = preprocess(raw)
    classification = model.predict(img)

    if classification == MY_CAT:
        open_door()

    machine.lightsleep(500)
```

